

bilgin_sherifov__customer_value_challenge

August 8, 2019

0.0.1 Submission for

0.1 *****: Data Challenge

by Bilgin Sherifov Date: ****--

Packages to install before running the notebook (with these you could run the cells till the "clustering2 session. For more advanced sessions you will need more packages):

1. Python 3.x
2. Pandas
3. Numpy
4. Matplotlib
5. Seaborn
6. Contextlib
7. xlrd

```
(python -m pip install --user numpy matplotlib pandas seaborn contextlib xlrd)
```

0.2 Note on running the notebook:

First, the notebook should be ran in the ~/code folder. Then, one should extract the tables in the data folder.

```
In [1]: # import numpy and pandas
import numpy as np
import pandas as pd

# import plotting packages
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import seaborn as sns; sns.set();
%matplotlib inline

# time-related libraries
import time as tm
import datetime as dt
from datetime import timedelta
```

```

# system and OS libraries
import os
from os.path import join as pj
import copy
import pickle
import gc; gc.collect()

#MSCL
from contextlib import contextmanager
import warnings
import itertools

from IPython.core.display import display, HTML
display(HTML("<style>.container { width:95% !important; }</style>"))

sns.set_context("talk") # paper, notebook, talk, poster

USE_DARK_THEME = False

<IPython.core.display.HTML object>

```

1 SOME CORE FUNCTIONS / CLASSES

```

In [2]: @contextmanager
def timer(title, n_blanks = 1):
    print(n_blanks * ' ' + 'STARTING: {}'.format(title))
    t0 = tm.time()
    yield
    print(n_blanks * ' ' + "FINISHED in {:.2f} min. \n".format((tm.time() - t0)/60))

In [3]: def to_datetime(df, date_columns):
    for date_column in date_columns:
        with timer(date_column, n_blanks=4):
            df[date_column] = pd.to_datetime(df[date_column], dayfirst=True, format='%m/%d/%Y')

    return df

```

2 WORKSPACE

2.1 Set paths and parameters

```

In [4]: if False:
    PATHS = {'root' : os.getcwd()}

    for new_path in ['data', 'results', 'WORKSPACE', 'info']:
        os.mkdir(pj(PATHS['root'], new_path))

```

```

        PATHS[new_path] = pj(PATHS['root'], new_path)
    else:
        #os.chdir("../")
        os.chdir(os.path.dirname(os.getcwd()))
        PATHS = {'root' : os.getcwd()}
        for new_path in ['data', 'results', 'WORKSPACE', 'info', 'code']:
            PATHS[new_path] = pj(PATHS['root'], new_path)
    [print(key,':', val) for key, val in PATHS.items()];

```

```

data : E:\JupyterNB\Job Applications\Shop Direct\data
code : E:\JupyterNB\Job Applications\Shop Direct\code
results : E:\JupyterNB\Job Applications\Shop Direct\results
info : E:\JupyterNB\Job Applications\Shop Direct\info
root : E:\JupyterNB\Job Applications\Shop Direct
WORKSPACE : E:\JupyterNB\Job Applications\Shop Direct\WORKSPACE

```

2.2 PLAN:

I'll define a customer as valuable if he/she is unlikely to default and also generates a lot of revenue. Therefore, a complete system should focus on both of these. However, due to time limitations, for this submission I'll focus only on the first part, i.e. whether a customer is likely to default or not.

Assessing the probability of default is basically a classification problem, while estimating the revenue that will be generated is a regression problem. Thus, in this submission I'll be focusing on classification. While there are many classifiers one can use, I'll use LightGBM, since it is fast, accurate, agnostic to data types (categorical or numeric), and has proven to be quite popular recently.

Before doing classification, though, I'll do basic data exploration with the purpose of selecting most relevant features. For this I'll use the fact that according to Bayes' theorem, the probability of defaulting or not given a set of features is proportional to the joint likelihood of observing these features given that a customer has defaulted or not, multiplied by the prior probability of defaulting or not (and normalised by the joint probability of observing the features).

It is difficult to calculate and visualize joint probabilities of features (conditional on default or not) of multidimensional data. Therefore, I'll take one feature at a time, akin Naive Bayes. If two features taken independently exhibit significantly different probability distributions for defaulted and non-defaulted visitor, one might safely assume that the joint probability of the two features will be even more distinct for the two classes of users.

I'll also use cluster analysis to assess whether I could separate customers into two distinct subsets on the basis of the features selected and I'll compare how the clusters overlap with whether a customer has defaulted or not. This, this will be one way of reassuring myself that I have selected a good set of features.

Thus, here is the plan:

1. Load the data and look at some basic things, like row and columns counts and data types etc
2. Compare visually probability distributions of various features, conditioned on whether the user defaulted or not.
3. Assess the quality of the features using some cluster analysis

4. Build a two-stage classification system, whereby first I build a classifier to detect customers that have very low probability of defaulting and filtering them out. Then, on the remaining set of data, which now should be much more class-balanced, I build a second classifier tht is tuned to better detect customers that are likely to default.
5. I also plot the features that are most important in the two stages of classification.

2.3 Quick exploration of the provided data

2.3.1 Read the tables and their legends and collect them into a dictionary

```
In [6]: with timer('Reading original data files'):
        dataset = {}
        for root, directories, files in os.walk(PATHS['data']):
            for file in files:
                if '.csv' in file:
                    try:
                        fileType = file.split('.')[0]
                        dataset[file.split('.')[0]] = {
                            'data': pd.read_csv(pj(PATHS['data'], file)),
                            'legend': pd.read_excel(pj(PATHS['info'], 'Data_Dictionary.xls'),
                                                    sheet_name=fileType)
                        }
                        print(4*' ' + 'Read:', file)
                    except:
                        print(4*' ' + "Couldn't read:", file)

        print(dataset.keys())

STARTING: Reading original data files
Read: Calendar.csv
Read: Customer.csv
Read: Order.csv
FINISHED in 0.06 min.

dict_keys(['Customer', 'Calendar', 'Order'])
```

2.3.2 Print table stats etc

```
In [7]: def table_stats(df, tableName):
        columns = list(df.columns)
        print('TABLE', tableName.upper())
        print(100 * '-')
        print('{0:<40s}: {1:,d}'.format('Number of rows', len(df)))
        print('{0:<40s}: {1:,d}'.format('Number of columns', len(columns)))
        print('\n{0:40s} {1:<30s}\n'.format('Column name', 'Number of unique values (data)'))
        for column in columns:
            print('{0:>40s}: {1:,d} ({2:})'.format(column, len(df[column].unique()), df[column].unique().tolist()))
        print(100 * '-')
```

Customers

Show table stats and a few rows

```
In [8]: tableName = 'Customer'
        table_stats(dataset[tableName]['data'], tableName)
        dataset[tableName]['data'].head(5)
```

TABLE CUSTOMER

```
-----
Number of rows           : 509,693
Number of columns        : 15
```

```
Column name              Number of unique values (data type)
```

```

        identifier: 509,693 (int64)
        title_desc: 11 (object)
principal_brand: 2 (object)
        birth_year: 96 (float64)
        credit_band: 25 (object)
        status_code: 3 (object)
        postcode_outward: 2,813 (object)
        date_start: 940 (object)
        date_default: 194 (object)
        date_completed: 1,002 (object)
        date_last_payment: 1,158 (object)
        open_to_buy_amt: 4,274 (float64)
LTIME_NET_SALES_AMT: 101,243 (float64)
        LTIME_NO_ORDERS: 167 (int64)
        LTIME_RETURNED_AMT: 23,665 (float64)
-----
```

```
Out [8]:  identifier title_desc principal_brand birth_year credit_band status_code \
0         37295      MISS          LAI      1972.0          Y1      active
1         315441      MRS          LEX      1966.0          Y1  completed
2         489894      MISS          LEX      1996.0          Y3      active
3         419603       MR          LEX      1946.0          Y1      active
4         444889      MISS          LEX      1968.0          U1      active

        postcode_outward date_start date_default date_completed date_last_payment \
0          NR30  19OCT2016          NaN          NaN      09MAR2017
1           MK5  08APR2016          NaN      21JAN2017      22APR2016
2          B 71  27AUG2016          NaN          NaN      13MAR2017
3          DE73  13DEC2016          NaN          NaN      14DEC2016
4          EX10  20JUN2014          NaN          NaN      02FEB2017

        open_to_buy_amt  LTIME_NET_SALES_AMT  LTIME_NO_ORDERS  LTIME_RETURNED_AMT
```

0	78.0	604.42	5	0.0
1	3000.0	34.39	1	0.0
2	4.0	512.94	5	0.0
3	1500.0	25.98	1	0.0
4	0.0	424.94	2	0.0

Show table legend

In [9]: dataset[tableName]['legend']

Out[9]:

	FIELD NAME	DATA TYPE	EXAMPLE \
0	identifier	integer	37295
1	title_desc	character	MISS
2	principal_brand	character	LAI
3	birth_year	integer	1972
4	credit_band	character	Y1
5	status_code	character	active
6	postcode_outward	character	NR30
7	date_start	date	2016-10-19 00:00:00
8	date_default	date	NaN
9	date_completed	date	NaN
10	date_last_payment	date	2017-03-09 00:00:00
11	open_to_buy_amt	decimal(8,2)	78
12	LTIME_NET_SALES_AMT	decimal(9,2)	604.42
13	LTIME_NO_ORDERS	integer	5
14	LTIME_RETURNED_AMT	decimal(9,2)	0

	DESCRIPTION
0	Individual identifier
1	Customer's title, such as Mr, Ms, Mrs, Dr etc.
2	Brand of customers account - usually matches t...
3	Year customer was born
4	Code that identifies the level of credit risk ...
5	status of account which can be - active - cust...
6	Outward for customers address
7	Date customer started trading
8	Date customer's account defaulted - blank if n...
9	Date customer's account was completed - blank ...
10	Date customer last made a payment on their acc...
11	Amount of credit still available on customer's...
12	Total Net Sales on the account
13	Total number of orders made on the account
14	Total Value of returned items on the account

Orders

Show table stats and a few rows

```
In [10]: tableName = 'Order'
         table_stats(dataset[tableName]['data'], tableName)
         dataset[tableName]['data'].head(5)
```

TABLE ORDER

```
-----
Number of rows                : 3,613,185
Number of columns            : 10
```

Column name Number of unique values (data type)

```

        identifier: 376,746 (int64)
        Brand: 2 (object)
    Account_Year_Week: 157 (int64)
        Week_Ending: 157 (object)
        Channel: 2 (object)
online_device_type_detail: 4 (object)
        Account_Type2: 2 (object)
    Gross_Demand_Pre: 9,305 (float64)
        New_Cust: 2 (object)
        Product_dept: 11 (object)
-----
```

```
Out [10]:
```

	identifier	Brand	Account_Year_Week	Week_Ending	Channel	\
0	37295	LAI	201652	23DEC2016	Online	
1	315441	LEX	201615	08APR2016	Online	
2	489894	LEX	201649	02DEC2016	Online	
3	419603	LEX	201651	16DEC2016	Online	
4	431206	LAI	201547	20NOV2015	Online	

	online_device_type_detail	Account_Type2	Gross_Demand_Pre	New_Cust	\
0	TABLET	Credit	27.00	Y	
1	DESKTOP	Credit	30.40	Y	
2	MOBILE	Credit	129.00	Y	
3	DESKTOP	Credit	21.99	Y	
4	MOBILE	Cash	29.75	Y	

	Product_dept
0	Dept G
1	Dept A
2	Dept B
3	Dept C
4	Dept D

Show table legend

```
In [11]: dataset[tableName]['legend']
```

```

Out[11]:
      FIELD NAME      DATA TYPE  EXAMPLE \
0      identifier      integer    37295
1              Brand      character    LAI
2      Account_Year_Week      integer    201652
3              Week_Ending      integer    1161223
4              Channel      character    Online
5      online_device_type_detail      character    TABLET
6              Account_Type2      character    Credit
7      Gross_Demand_Pre_Credit      decimal(9.2)    27
8              New_Cust      character    Y
9              Product_dept      character    DEPT A

      DESCRIPTION
0      Individual identifier
1      Identifies the Brand, LEX = VERY. LAI = Little...
2              Week and Year
3      Date on last day of the week, i.e. the Friday
4      Either online purchase or offline purchase
5              Device used for purchase
6      Whether the account a cash account or a credit...
7              Price customer pays for order
8      Y = new customer, N = existing customer
9              Department name

```

Calendar

Show table stats and a few rows

```

In [12]: tableName = 'Calendar'
         table_stats(dataset[tableName]['data'], tableName)
         dataset[tableName]['data'].head(5)

```

TABLE CALENDAR

```

-----
Number of rows      : 1,099
Number of columns   : 20

```

Column name Number of unique values (data type)

```

      ACCOUNT_PERIOD: 12 (int64)
      ACCOUNT_PERIOD_WEEK: 6 (int64)
      ACCOUNT_WEEK: 53 (int64)
      ACCOUNT_YEAR_PERIOD: 36 (object)
      ACCOUNT_YEAR_PERIOD_REL_NO: 36 (int64)
      ACCOUNT_YEAR_WEEK: 157 (object)
      ACCOUNT_YEAR_WEEK_REL_NO: 157 (int64)
      ACCOUNT_YEAR: 3 (object)
      ACCOUNT_YEAR_REL_NO: 3 (int64)

```


CAL_DATE: 1,099 (object)
 CAL_DATE_REL_NO: 1,099 (object)
 CAL_DAY_ID: 7 (int64)
 CAL_WEEK: 53 (int64)
 CAL_PERIOD: 13 (int64)
 CAL_PERIOD_WEEK: 5 (int64)
 CAL_YEAR_MONTH: 37 (object)
 DAY_ID: 6 (int64)
 SEASON: 6 (object)
 SEASON_REL_NO: 6 (int64)
 SEASON_WEEK: 27 (int64)

Out [12]:

	ACCOUNT_PERIOD	ACCOUNT_PERIOD_WEEK	ACCOUNT_WEEK	ACCOUNT_YEAR_PERIOD	\
0	2	4	8	201,502	
1	11	2	46	201,611	
2	5	2	19	201,605	
3	8	4	34	201,408	
4	1	1	1	201,501	

	ACCOUNT_YEAR_PERIOD_REL_NO	ACCOUNT_YEAR_WEEK	ACCOUNT_YEAR_WEEK_REL_NO	\
0	-25	201,508	-108	
1	-4	201,646	-18	
2	-10	201,619	-45	
3	-31	201,434	-134	
4	-26	201,501	-115	

	ACCOUNT_YEAR	ACCOUNT_YEAR_REL_NO	CAL_DATE	CAL_DATE_REL_NO	CAL_DAY_ID	\
0	2,015	-2	18/02/2015	-758	4	
1	2,016	-1	10/11/2016	-127	5	
2	2,016	-1	30/04/2016	-321	0	
3	2,014	-3	21/08/2014	-939	5	
4	2,015	-2	31/12/2014	-807	4	

	CAL_WEEK	CAL_PERIOD	CAL_PERIOD_WEEK	CAL_YEAR_MONTH	DAY_ID	SEASON	\
0	8	2	4	201,502	4	SS2015	
1	45	12	1	201,611	5	AW2016	
2	17	5	1	201,604	1	SS2016	
3	34	9	2	201,408	5	AW2014	
4	1	1	1	201,412	4	SS2015	

	SEASON_REL_NO	SEASON_WEEK
0	-4	8
1	-1	19
2	-2	19
3	-5	8
4	-4	1

Show table legend

In [13]: dataset[tableName]['legend']

Out [13]:

	FIELD NAME	DATA TYPE	EXAMPLE \
0	ACCOUNT_PERIOD	byteint	2
1	ACCOUNT_PERIOD_WEEK	byteint	4
2	ACCOUNT_WEEK	byteint	8
3	ACCOUNT_YEAR_PERIOD	integer	201502
4	ACCOUNT_YEAR_PERIOD_REL_NO	integer	-25
5	ACCOUNT_YEAR_WEEK	integer	201508
6	ACCOUNT_YEAR_WEEK_REL_NO	integer	-108
7	ACCOUNT_YEAR	smallint	2015
8	ACCOUNT_YEAR_REL_NO	integer	-2
9	CAL_DATE	date	2015-02-18 00:00:00
10	CAL_DATE_REL_NO	integer	-758
11	CAL_DAY_ID	byteint	4
12	CAL_WEEK	byteint	8
13	CAL_PERIOD	byteint	2
14	CAL_PERIOD_WEEK	byteint	4
15	CAL_YEAR_MONTH	integer	201502
16	DAY_ID	byteint	4
17	SEASON	character(6)	SS2015
18	SEASON_REL_NO	integer	-4
19	SEASON_WEEK	byteint	8

	DESCRIPTION
0	Accounting period. Values from 1 to 12.
1	Accounting period week. Values from 1 to 6.
2	Accounting week. Values from 1 to 53.
3	Accounting year period. Format yyy,ypp where y...
4	Accounting year period relative number. It re...
5	Accounting year week. Format yyy,yww where yyy...
6	Accounting year week relative number. It repre...
7	Accounting year.
8	Accounting year relative number. It represents...
9	Calendar date.
10	Calendar date relative number. Number that rep...
11	Calendar day ID. From 0 to 6: (Sat 0, Sun 1, M...
12	Calendar week. One calendar year has up to 53 ...
13	Calendar period. Values from 1 to 53.
14	Calendar period week. From 1 up to 5
15	Calendar year & month. Format yyy,ypp
16	Calendar day ID. From 1 to 6: (Sat 1, Sun 1, M...
17	Season name. Format: 'SSyyyy' or 'AWyyyy' wher...
18	Season relative number. It represents the numb...
19	Week number within each season.

2.4 A little more indepth exploration of the provided data

2.4.1 Customers table alone

Data augmenttion and transformation At this stage, I'll look at the probability distributions of certain fields as a function of whether the customer has defaulted or not and see if they are different for each customer type.

1. First, I'll append a column, dafaulted, which will be a flag for whether customer deefaulted or not (True if defaulted, False otherwise)
2. Next, I'll convert all date-related columns to datetime data format
3. Next, I'll add a few more date-related columns

```
In [14]: with timer('Appending a column that is a flag for whether customer deefaulted or not'):
        # make a copy of the table
        dfCustomer = copy.deepcopy(dataset['Customer']['data'])
        # ad a columns that indictes defaulted customer or not and set all to False at fi
        dfCustomer['defaulted'] = False
        # get row index of defaulted customers
        idxDefaulted = dfCustomer[dfCustomer['date_default'].notnull()].index
        # set the defaulted flag to True for those rows
        dfCustomer.loc[idxDefaulted, 'defaulted'] = True

        with timer('Converting date columns to datetime'):
            dateColumns = ['date_start', 'date_default', 'date_completed', 'date_last_payment']
            dfCustomer = to_datetime(dfCustomer, dateColumns)

        with timer('Appending date-related columns'):
            # get year and month of start date
            dfCustomer['year_start'] = dfCustomer['date_start'].dt.year
            dfCustomer['month_start'] = dfCustomer['date_start'].dt.month
            # get year and month of last payment date
            dfCustomer['year_last_payment'] = dfCustomer['date_last_payment'].dt.year
            dfCustomer['month_last_payment'] = dfCustomer['date_last_payment'].dt.month
            # calculate time in days between last payment date and account creation date
            dfCustomer['days_between_last_payment_and_start'] = (dfCustomer['date_last_payment'] - dfCustomer['date_start']).dt.days
            # set identifier as a row index
            dfCustomer = dfCustomer.set_index('identifier', drop=False)

        dfCustomer.head(5)
```

STARTING: Appending a column that is a flag for whether customer deefaulted or not
FINISHED in 0.00 min.

STARTING: Converting date columns to datetime
STARTING: date_start
FINISHED in 0.02 min.

STARTING: date_default
FINISHED in 0.00 min.

STARTING: date_completed
FINISHED in 0.01 min.

STARTING: date_last_payment
FINISHED in 0.01 min.

FINISHED in 0.04 min.

STARTING: Appending date-related columns
FINISHED in 0.00 min.

```
Out[14]:
```

identifier	title_desc	principal_brand	birth_year	credit_band	\
37295	37295	MISS	LAI	1972.0	Y1
315441	315441	MRS	LEX	1966.0	Y1
489894	489894	MISS	LEX	1996.0	Y3
419603	419603	MR	LEX	1946.0	Y1
444889	444889	MISS	LEX	1968.0	U1

identifier	status_code	postcode_outward	date_start	date_default	\
37295	active	NR30	2016-10-19	NaT	
315441	completed	MK5	2016-04-08	NaT	
489894	active	B 71	2016-08-27	NaT	
419603	active	DE73	2016-12-13	NaT	
444889	active	EX10	2014-06-20	NaT	

identifier	date_completed	...	\
37295	NaT	...	
315441	2017-01-21	...	
489894	NaT	...	
419603	NaT	...	
444889	NaT	...	

identifier	open_to_buy_amt	LTIME_NET_SALES_AMT	LTIME_NO_ORDERS	\
37295	78.0	604.42	5	
315441	3000.0	34.39	1	
489894	4.0	512.94	5	
419603	1500.0	25.98	1	
444889	0.0	424.94	2	

	LTIME_RETURNED_AMT	defaulted	year_start	month_start	\
identifier					
37295	0.0	False	2016	10	
315441	0.0	False	2016	4	
489894	0.0	False	2016	8	
419603	0.0	False	2016	12	
444889	0.0	False	2014	6	

	year_last_payment	month_last_payment	\
identifier			
37295	2017.0	3.0	
315441	2016.0	4.0	
489894	2017.0	3.0	
419603	2016.0	12.0	
444889	2017.0	2.0	

	days_between_last_payment_and_start
identifier	
37295	141.0
315441	14.0
489894	198.0
419603	1.0
444889	958.0

[5 rows x 21 columns]

Plot distributions for numeric features

```
In [15]: def plot_prob_numeric_feature_given_defaulted(data, features, log_transformed_fields =
n_rows = len(features)
plt.close('all')
plt.figure(figsize=(24, n_rows * 4))
for j, feature in enumerate(features):
    plt.subplot(n_rows, 2, j*2+1)
    df = copy.deepcopy(data[data['principal_brand'] == 'LAI'])
    if feature in log_transformed_fields:
        # get two index of non-negative values only, because I'll be taking the log
        idx_non_negative = df[df[feature] >= 0].index
        df = df.loc[idx_non_negative]
        series_defaulted = (df[df['defaulted']][feature]+.001).apply(np.log10)
        series_non_defaulted = (df[~df['defaulted']][feature]+.001).apply(np.log10)
    else:
        idx_non_null = df[df[feature].notnull()].index
        df = df.loc[idx_non_null]
        series_defaulted = df[df['defaulted']][feature]
        series_non_defaulted = df[~df['defaulted']][feature]

    # plot normed histogram of log_value for defaulted customers
```

```

sns.distplot(
    series_defaulted,
    hist = True,
    label='defaulted')
# plot normed histogram of log_value for non-defaulted customers
sns.distplot(
    series_non_defaulted,
    hist = True,
    label='non-defaulted')

if feature in log_transformed_fields:
    plt.xlabel(feature + ' (transformed to $\log_{10}$)')
plt.legend()
if j == 0:
    plt.title('Histograms for Littlewoods')

plt.subplot(n_rows, 2, j*2+2)
df = copy.deepcopy(data[data['principal_brand'] == 'LEX'])
if feature in log_transformed_fields:
    # get two index of non-negative values only, because I'll be taking the log
    idx_non_negative = df[df[feature] >= 0].index
    df = df.loc[idx_non_negative]
    series_defaulted = (df[df['defaulted']][feature]+.001).apply(np.log10)
    series_non_defaulted = (df[~df['defaulted']][feature]+.001).apply(np.log10)
else:
    idx_non_null = df[df[feature].notnull()].index
    df = df.loc[idx_non_null]
    series_defaulted = df[df['defaulted']][feature]
    series_non_defaulted = df[~df['defaulted']][feature]

# plot normed histogram of log_value for defaulted customers
sns.distplot(
    series_defaulted,
    hist = True,
    label='defaulted')
# plot normed histogram of log_value for non-defaulted customers
sns.distplot(
    series_non_defaulted,
    hist = True,
    label='non-defaulted')

if feature in log_transformed_fields:
    plt.xlabel(feature + ' (transformed to $\log_{10}$)')
plt.legend()
if j == 0:
    plt.title('Histograms for VERY')

```

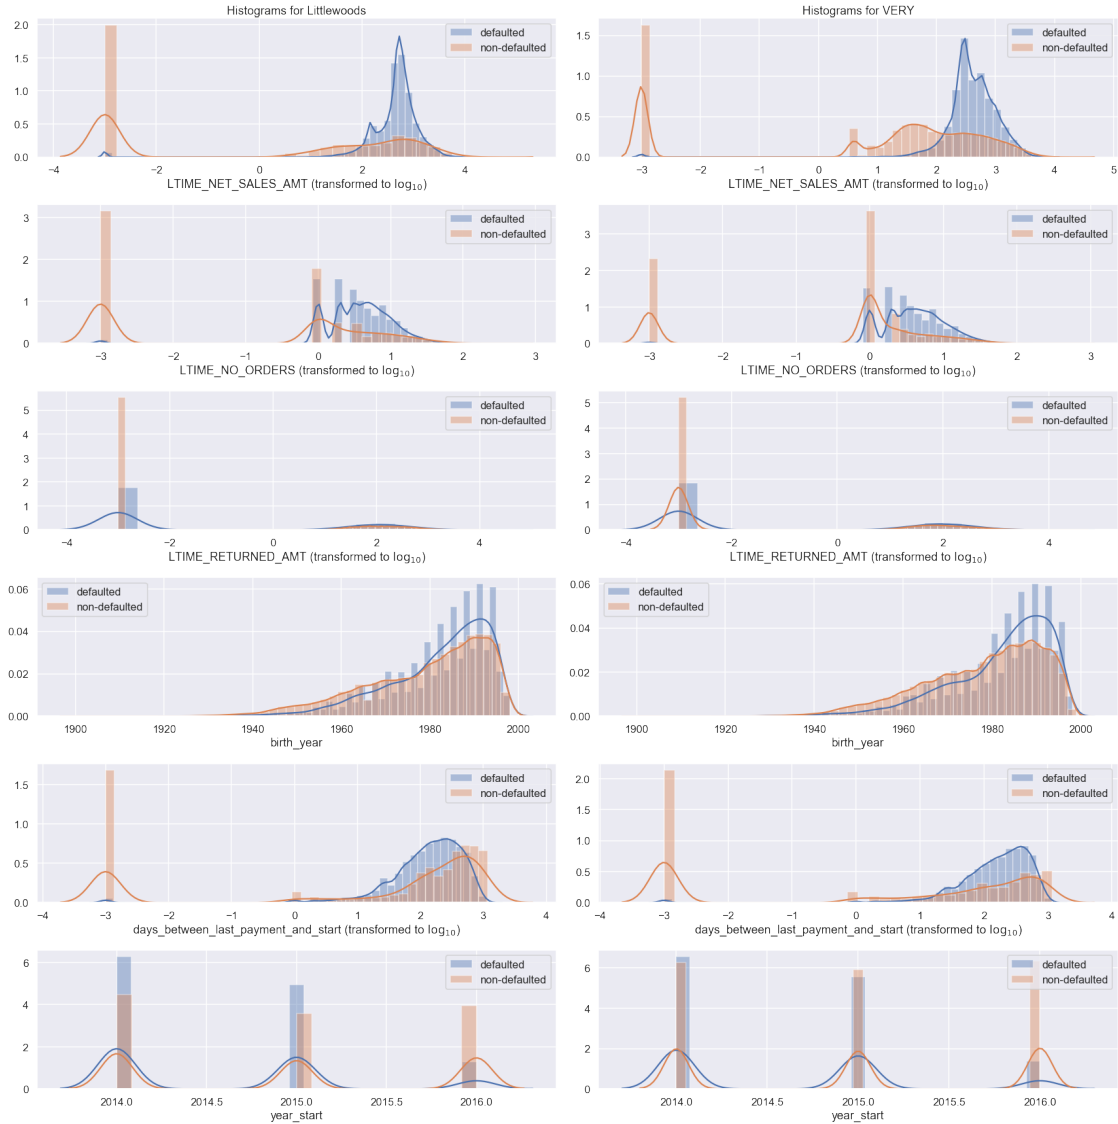
```

plt.tight_layout()

In [16]: plot_prob_numeric_feature_given_defaulted(
    dfCustomer,
    [
        'LTIME_NET_SALES_AMT',
        'LTIME_NO_ORDERS',
        'LTIME_RETURNED_AMT',
        'birth_year',
        'days_between_last_payment_and_start',
        'year_start'
    ],
    log_transformed_fields = [
        'LTIME_NET_SALES_AMT',
        'LTIME_NO_ORDERS',
        'LTIME_RETURNED_AMT',
        'days_between_last_payment_and_start'
    ]
)

C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```



Observations:

- **Total Net Sales:** In both brand accounts, there is a difference in the distribution of Total Net Sales for defaulted and non-defaulted users. For defaulted users this distribution peaks at higher values and is more concentrated. Thus, this feature should be included in any model that tries to predict the chance of a user defaulting. The difference in distributions between defaulted and non-defaulted users seems to be higher for "VERY"
- **Total number of orders:** In both brand accounts, there is a difference in the distribution of Total number of orders for defaulted and non-defaulted users. For defaulted users this distribution is skewed towards higher values. Thus, this feature should be included in any model that tries to predict the chance of a user

defaulting.

There seems to be no visible difference in distributions between defaulted and non-defaulted users across brands.

- Total Value of returned items: In both brand accounts, there is NO visible difference in the distribution of Total Value of returned items for defaulted and non-defaulted users.

Thus, this feature could be excluded in any model that tries to predict the chance of a user defaulting.

- Year customer was born: In both brand accounts, the distribution Year customer was born is skewed towards higher values, implying younger customers predominate. However, it looks like the distribution is more concentrated around young customers for defaulted users. So, there is a difference in the distribution of Year customer was born for defaulted and non-defaulted users.

Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.

There seems to be no difference in distributions between defaulted and non-defaulted users across brands.

- Days between last payment and start date: In both brand accounts, there is a difference in the distribution of Days between last payment and start date for defaulted and non-defaulted users. For defaulted users this distribution is centered and peaked towards lower values.

Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.

There seems to be no visible difference in distributions between defaulted and non-defaulted users across brands.

- Year account started: In both brand accounts, there is a difference in the distribution of Year account started for defaulted and non-defaulted users. For defaulted users this distribution is skewed towards earlier years. This is probably expected, as the longer the account has been active, probably the more chances to default.

Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.

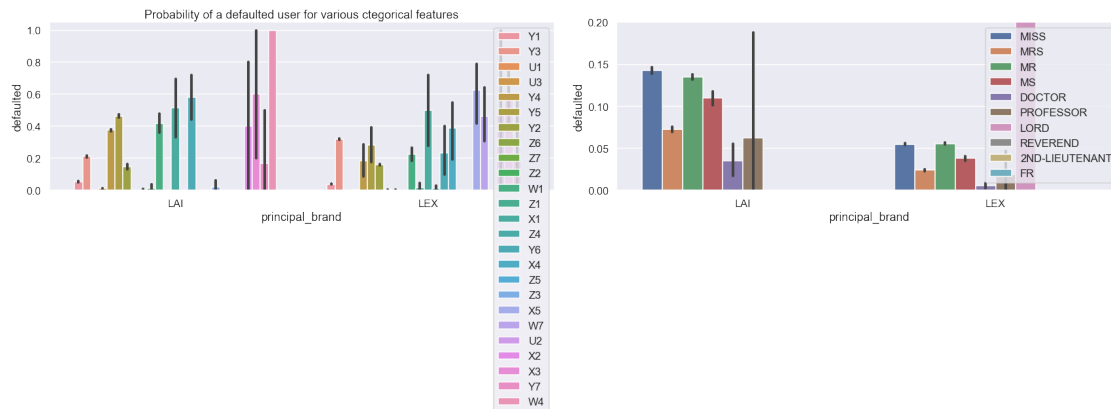
The difference in distributions between defaulted and non-defaulted users seems to be higher for "Littlewoods"

Plot distributions for categorical features

```
In [17]: def plot_prob_categorical_feature_given_defaulted(data, features):
    n_rows = len(features)
    plt.close('all')
    plt.figure(figsize=(24, n_rows * 6))
    for j, feature in enumerate(features):
        plt.subplot(n_rows, 2, j+1)
        sns.barplot(x="principal_brand", y="defaulted", hue=feature, estimator = np.mean)
        if j == 0:
            plt.title('Probability of a defaulted user for various categorical features')
        if feature == 'title_desc':
            plt.ylim([0, .2])
        plt.legend(loc='upper right')
    plt.tight_layout()
```

```
In [18]: plot_prob_categorical_feature_given_defaulted(dfCustomer, ['credit_band', 'title_desc'])
```

```
C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



Observations:

- Level of credit risk: In both brand accounts, there is a difference in the distribution of Level of credit risk for defaulted and non-defaulted users. This is not surprising, since these probably were already based on prior credit risk assesment. Nevertheless, this feature will be included in later stages.
Also, there seems to be some difference in the distributions between defaulted and non-defaulted users accross brands.
- Customer's title: In both brand accounts, there seems to have a statistically significant difference in the distribution of Customer's title for defaulted and non-defaulted users. For defaulted users this distribution peaks at higher values and is more concentrated.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
There seems to be no difference in distributions between defaulted and non-defaulted users accross brands (excluding the title LORD, which has one only sample per group).

Select only relevant features and one-hot-encode the selected categorical features

```
In [19]: with timer('Removing un-informative columns:'):
        selected_columns_customer = ['identifier', 'defaulted']

        selected_numeric_columns_customer = [
            'LTIME_NET_SALES_AMT',
            'LTIME_NO_ORDERS',
            'birth_year',
            'days_between_last_payment_and_start',
            'year_start'
        ]
        selected_columns_customer.extend(selected_numeric_columns_customer)
```

```

dfCustomer_1 = copy.deepcopy(dfCustomer[selected_columns_customer])

with timer('One-hot-encoding credit_band', n_blanks=4):
    # one-hot-encoding by credit_band and append to previous
    dfTemp = dfCustomer.groupby(['identifier', 'credit_band']).size().unstack(1).fillna(0)
    dfTemp.columns = ['credit_band.' + clmn for clmn in list(dfTemp)]
    dfCustomer_1 = dfCustomer_1.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

with timer('One-hot-encoding title_desc', n_blanks=4):
    # one-hot-encoding by title_description and append to previous
    dfTemp = dfCustomer.groupby(['identifier', 'title_desc']).size().unstack(1).fillna(0)
    dfTemp.columns = ['title.' + clmn for clmn in list(dfTemp)]
    dfCustomer_1 = dfCustomer_1.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

# free memory by deleting the dfCustomer table
del dfCustomer
gc.collect();

dfCustomer_1.head(5)

```

STARTING: Removing un-informative columns:
FINISHED in 0.00 min.

STARTING: One-hot-encoding credit_band

C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\ipykernel__main__.py:16: FutureWarning: Defaulting to column, but this will raise an ambiguity error in a future version

FINISHED in 0.02 min.

STARTING: One-hot-encoding title_desc

C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\ipykernel__main__.py:22: FutureWarning: Defaulting to column, but this will raise an ambiguity error in a future version

FINISHED in 0.02 min.

```

Out[19]:

```

	identifier	defaulted	LTIME_NET_SALES_AMT	LTIME_NO_ORDERS	\
identifier					
37295	37295	False	604.42	5	
315441	315441	False	34.39	1	
489894	489894	False	512.94	5	
419603	419603	False	25.98	1	

444889	444889	False	424.94	2
--------	--------	-------	--------	---

	birth_year	days_between_last_payment_and_start	year_start	\
identifier				
37295	1972.0		141.0	2016
315441	1966.0		14.0	2016
489894	1996.0		198.0	2016
419603	1946.0		1.0	2016
444889	1968.0		958.0	2014

	credit_band.U1	credit_band.U2	credit_band.U3	...	\
identifier				...	
37295	0	0	0	...	
315441	0	0	0	...	
489894	0	0	0	...	
419603	0	0	0	...	
444889	1	0	0	...	

	title.2ND-LIEUTENANT	title.DOCTOR	title.FR	title.LORD	\
identifier					
37295	0	0	0	0	
315441	0	0	0	0	
489894	0	0	0	0	
419603	0	0	0	0	
444889	0	0	0	0	

	title.MISS	title.MR	title.MRS	title.MS	title.PROFESSOR	\
identifier						
37295	1	0	0	0	0	
315441	0	0	1	0	0	
489894	1	0	0	0	0	
419603	0	1	0	0	0	
444889	1	0	0	0	0	

	title.REVEREND
identifier	
37295	0
315441	0
489894	0
419603	0
444889	0

[5 rows x 42 columns]

2.4.2 Include Orders table, as well

Data augmentation and transformation

1. First, append the 'identifier' and 'defaulted' columns from the Customer table into the

Order table. Merge on 'identifier'.

2. Then convert the data type of the 'Week_Ending' column into datetime.
3. Then, append more columns related to purchase date and counts
4. Finally, transform the orders table into a table indexed by row and columns are either one-hot-encoded version of the original categorical features or are including stats of the numeric features. I'll call it dfCustomers_2

```
In [20]: with timer('Appending defaulted column from Customers table'):
        # make a copy of the table
        dfOrder = copy.deepcopy(dataset['Order']['data'])
        # add a column that indicates defaulted customer or not and set all to False at first
        dfOrder = dfOrder.merge(dfCustomer_1[['identifier', 'defaulted']], how = 'left',
                                left_index=True, right_index=True)

        with timer('Converting date columns to datetime'):
            dateColumns = ['Week_Ending']
            dfOrder = to_datetime(dfOrder, dateColumns)

        with timer('Appending columns related to purchase date and counts'):
            # Create a grouped object (by identifier)
            grouped = dfOrder.groupby('identifier')
            # attach a column indicating days between last and first order
            dfOrder = dfOrder.merge(
                pd.DataFrame({'days_between_last_and_first_order':
                             (grouped['Week_Ending'].max()-grouped['Week_Ending'].min())
                             ,
                             'left_on':'identifier',
                             'right_on':'identifier'}),
                left_index=True, right_index=True)
            # attach a column indicating number of orders
            dfOrder = dfOrder.merge(
                pd.DataFrame({'number_of_orders':grouped.size()}),
                left_index=True, right_index=True,
                left_on='identifier', right_on='identifier')

        dfOrder.head(5)
```

STARTING: Appending defaulted column from Customers table

C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\ipykernel__main__.py:6: FutureWarning: Defaulting to column, but this will raise an ambiguity error in a future version

FINISHED in 0.03 min.

STARTING: Converting date columns to datetime

STARTING: Week_Ending

FINISHED in 0.13 min.

FINISHED in 0.13 min.

STARTING: Appending columns related to purchase date and counts

FINISHED in 0.05 min.

```
Out[20]:
```

	identifier	Brand	Account_Year_Week	Week_Ending	Channel	\
0	37295	LAI	201652	2016-12-23	Online	
1	37295	LAI	201643	2016-10-21	Online	
2	37295	LAI	201651	2016-12-16	Online	
3	315441	LEX	201615	2016-04-08	Online	
4	489894	LEX	201649	2016-12-02	Online	

	online_device_type_detail	Account_Type2	Gross_Demand_Pre	New_Cust	\
0	TABLET	Credit	27.00	Y	
1	DESKTOP	Credit	249.99	Y	
2	TABLET	Credit	24.00	Y	
3	DESKTOP	Credit	30.40	Y	
4	MOBILE	Credit	129.00	Y	

	Product_dept	defaulted	days_between_last_and_first_order	number_of_orders
0	Dept G	False	63	3
1	Dept C	False	63	3
2	Dept B	False	63	3
3	Dept A	False	0	1
4	Dept B	False	91	4

```
In [21]: with timer('Transforming Orders table into Customer vs Features table'):
    with timer('One-hot-encoding Brand', n_blanks=4):
        # one-hot-encoding by Brand, row index by identifier
        dfCustomer_2 = dfOrder.groupby(['identifier', 'Brand']).size().unstack(1).fillna(0)
        dfCustomer_2.columns = ['Brand.' + clmn for clmn in list(dfCustomer_2)]

    with timer('One-hot-encoding Channel', n_blanks=4):
        # one-hot-encoding by Channel and append to previous
        dfTemp = dfOrder.groupby(['identifier', 'Channel']).size().unstack(1).fillna(0)
        dfTemp.columns = ['Channel.' + clmn for clmn in list(dfTemp)]
        dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

    with timer('One-hot-encoding Device Type', n_blanks=4):
        # one-hot-encoding by online_device_type_detail and append to previous
        dfTemp = dfOrder.groupby(['identifier', 'online_device_type_detail']).size().unstack(1).fillna(0)
        dfTemp.columns = ['device.' + clmn for clmn in list(dfTemp)]
        dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

    with timer('One-hot-encoding Account Type', n_blanks=4):
        # one-hot-encoding by online_device_type_detail and append to previous
```

```

dfTemp = dfOrder.groupby(['identifier', 'Account_Type2']).size().unstack(1).fillna(0)
dfTemp.columns = ['account_type.' + clmn for clmn in list(dfTemp)]
dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

with timer('One-hot-encoding Customer Type', n_blanks=4):
    # rename Y to New and N to Returning
    idx = dfOrder[dfOrder['New_Cust'] == 'Y'].index
    dfOrder.loc[idx, 'New_Cust'] = 'new'
    idx = dfOrder[dfOrder['New_Cust'] == 'N'].index
    dfOrder.loc[idx, 'New_Cust'] = 'returning'
    # one-hot-encoding by online_device_type_detail and append to previous
    dfTemp = dfOrder.groupby(['identifier', 'New_Cust']).size().unstack(1).fillna(0)
    dfTemp.columns = ['customer_type.' + clmn for clmn in list(dfTemp)]
    dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

with timer('One-hot-encoding Product Department', n_blanks=4):
    # one-hot-encoding by online_device_type_detail and append to previous
    dfTemp = dfOrder.groupby(['identifier', 'Product_dept']).size().unstack(1).fillna(0)
    dfTemp.columns = ['department.' + clmn for clmn in list(dfTemp)]
    dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

with timer('Adding stats columns for Gross Demand Pre Credit', n_blanks=4):
    # one-hot-encoding by online_device_type_detail and append to previous
    grouped = dfOrder.groupby('identifier')
    dfTemp = grouped['Gross_Demand_Pre'].agg(['min', 'max', 'mean', 'median', 'std'])
    dfTemp.columns = ['Gross_Demand_Pre_Credit.' + clmn for clmn in list(dfTemp)]
    dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)
    dfCustomer_2['Gross_Demand_Pre_Credit.std'].fillna(0, inplace=True),

with timer('Adding stats columns for number_of_orders', n_blanks=4):
    # one-hot-encoding by online_device_type_detail and append to previous
    dfTemp = grouped[['number_of_orders']].mean()
    dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)

with timer('Adding stats columns for Days Between Last and First Order', n_blanks=4):
    # one-hot-encoding by online_device_type_detail and append to previous
    dfTemp = grouped[['days_between_last_and_first_order']].mean()
    dfCustomer_2 = dfCustomer_2.merge(dfTemp, how = 'inner', left_index=True, right_index=True)
    dfCustomer_2['avg_days_between_orders'] = (
        dfCustomer_2['days_between_last_and_first_order']
        / dfCustomer_2['number_of_orders'])

with timer('Appending the defaulted flag column', n_blanks=4):
    dfCustomer_2 = dfCustomer_2.merge(dfCustomer_1[['defaulted']], how = 'inner', left_index=True, right_index=True)

dfCustomer_2.head(5)

```

STARTING: Transforming Orders table into Customer vs Features table

STARTING: One-hot-encoding Brand

FINISHED in 0.02 min.

STARTING: One-hot-encoding Channel

FINISHED in 0.01 min.

STARTING: One-hot-encoding Device Type

FINISHED in 0.01 min.

STARTING: One-hot-encoding Account Type

FINISHED in 0.01 min.

STARTING: One-hot-encoding Customer Type

FINISHED in 0.04 min.

STARTING: One-hot-encoding Product Department

FINISHED in 0.02 min.

STARTING: Adding stats columns for Gross Demand Pre Credit

FINISHED in 0.01 min.

STARTING: Adding stats columns for number_of_orders

FINISHED in 0.00 min.

STARTING: Adding stats columns for Days Between Last_and First Order

FINISHED in 0.00 min.

STARTING: Appending the defaulted flag column

FINISHED in 0.00 min.

FINISHED in 0.13 min.

```
Out[21]:
```

	Brand.LAI	Brand.LEX	Channel.Offline	Channel.Online	\
identifier					
1	0	31	0	31	
2	0	1	0	1	
3	0	2	0	2	
4	0	1	0	1	
5	0	9	0	9	

	device.DESKTOP	device.MOBILE	device.TABLET	account_type.Cash	\
identifier					
1	0	27	4	0	
2	0	1	0	1	
3	0	2	0	0	

4	1	0	0	0
5	1	3	5	0

	account_type.Credit	customer_type.new	...	\
identifier				
1	31	2	...	
2	0	1	...	
3	2	1	...	
4	1	1	...	
5	9	7	...	

	department.Dept X	Gross_Demand_Pre_Credit.min	\
identifier			
1	0	7.00	
2	0	79.99	
3	0	59.99	
4	0	79.99	
5	0	9.90	

	Gross_Demand_Pre_Credit.max	Gross_Demand_Pre_Credit.mean	\
identifier			
1	110.00	38.137097	
2	79.99	79.990000	
3	199.00	129.495000	
4	79.99	79.990000	
5	409.00	210.100000	

	Gross_Demand_Pre_Credit.median	Gross_Demand_Pre_Credit.std	\
identifier			
1	35.000	21.695798	
2	79.990	0.000000	
3	129.495	98.294914	
4	79.990	0.000000	
5	215.000	194.573707	

	number_of_orders	days_between_last_and_first_order	\
identifier			
1	31	532	
2	1	0	
3	2	70	
4	1	0	
5	9	175	

	avg_days_between_orders	defaulted
identifier		
1	17.161290	False
2	0.000000	False
3	35.000000	False

4	0.000000	False
5	19.444444	False

[5 rows x 31 columns]

Plot distributions for numeric features

```
In [22]: def plot_prob_numeric_feature_given_defaulted_orders(data, features, log_transformed_
n_rows = len(features)
plt.close('all')
plt.figure(figsize=(24, n_rows * 4))
for j, feature in enumerate(features):
    plt.subplot(n_rows, 2, j+1)
    df = copy.deepcopy(data)
    if feature in log_transformed_fields:
        series_defaulted = (df[df['defaulted']][feature]+.001).apply(np.log10)
        series_non_defaulted = (df[~df['defaulted']][feature]+.001).apply(np.log10)
    else:
        series_defaulted = df[df['defaulted']][feature]
        series_non_defaulted = df[~df['defaulted']][feature]

    # plot normed histogram of log_value for defaulted customers
    sns.distplot(
        series_defaulted,
        hist = True,
        label='defaulted')
    # plot normed histogram of log_value for non-defaulted customers
    sns.distplot(
        series_non_defaulted,
        hist = True,
        label='non-defaulted')

    if feature in log_transformed_fields:
        plt.xlabel(feature + ' (transformed to $\log_{10}$)')
    plt.legend()
    if j == 0:
        plt.title('Histograms')

plt.tight_layout()

In [23]: plot_prob_numeric_feature_given_defaulted_orders(
dfCustomer_2,
[
    'Gross_Demand_Pre_Credit.min',
    'Gross_Demand_Pre_Credit.max',
    'Gross_Demand_Pre_Credit.mean',
    'Gross_Demand_Pre_Credit.median',
```

```

        'Gross_Demand_Pre_Credit.std',
        'days_between_last_and_first_order',
        'number_of_orders',
        'avg_days_between_orders'
    ],
    log_transformed_fields = [
        'Gross_Demand_Pre_Credit.min',
        'Gross_Demand_Pre_Credit.max',
        'Gross_Demand_Pre_Credit.mean',
        'Gross_Demand_Pre_Credit.median',
        'Gross_Demand_Pre_Credit.std',
        'days_between_last_and_first_order',
        'number_of_orders',
        'avg_days_between_orders'
    ]
)

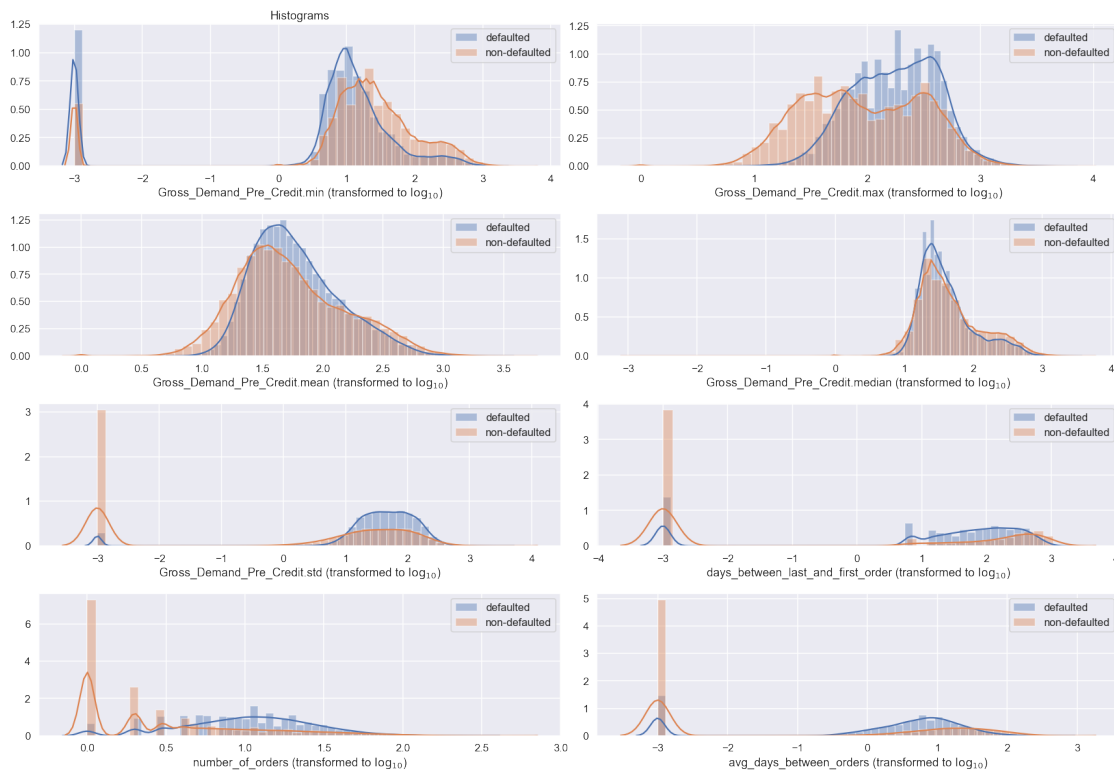
```

C:\Users\bilgin\Anaconda3\envs\ml_adv_py35\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: np.add.reduce is deprecated. Use np.add.reduce instead.

```

return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```



Observations:

- Minimum value of Gross_Demand_Pre_Credit: There is a difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is skewed towards lower values.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Maximum value of Gross_Demand_Pre_Credit: There is a difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is skewed towards higher values.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Mean value Gross_Demand_Pre_Credit: There is a slight difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is peaked at higher values.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Median value of Gross_Demand_Pre_Credit: There is not a visible difference in the distribution of this feature for defaulted and non-defaulted users.
Thus, this feature could be excluded from any model that tries to predict the chance of a user defaulting.
- Standard Deviation in Gross_Demand_Pre_Credit: There is not a visible difference in the distribution of this feature for defaulted and non-defaulted users, except for the very high peak at zero for non-defaulted users, due to customers with single purchase. Thus, this feature could be excluded from any model that tries to predict the chance of a user defaulting.
- Days between last payment and first purchase: There is a difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is skewed towards lower values.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Number of Orders: There is a difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is skewed towards higher values. Also, there are many non-defaulted customers with a single purchase.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Average Days Between Orders: There is a difference in the distribution of this feature for defaulted and non-defaulted users. For defaulted users the distribution is skewed towards higher values. Also, there are many non-defaulted customers with a single purchase.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.

Plot distributions for categorical features

```
In [24]: def plot_prob_categorical_unstacked_features_given_defaulted(df, features):
            idx_defaulted = df[df['defaulted']].index
            idx_non_defaulted = df[~df['defaulted']].index

            all_columns = list(df.columns)
            selected_columns = []
```

```

cnt_rows = len(features)
for feature in features:
    feature_columns = [column for column in all_columns if column.split('.')[0] ==
                        selected_columns.append(feature_columns)

plt.close('all')
plt.figure(figsize=(16, 4 * cnt_rows))
for j, columns in enumerate(selected_columns):
    df_temp = df[columns].divide(df['number_of_orders'],axis=0).merge(df[['defaulted']])
    plt.subplot(cnt_rows,2,j*2 + 1)
    df_probs = df_temp.loc[idx_defaulted][columns].mean()
    df_probs /= df_probs.sum()
    df_probs.plot.bar()
    plt.ylim([0,1])
    plt.title(features[j] + ': Defaulted')

    plt.subplot(cnt_rows,2,j*2 + 2)
    df_probs = df_temp.loc[idx_non_defaulted][columns].mean()
    df_probs /= df_probs.sum()
    df_probs.plot.bar()
    plt.ylim([0,1])
    plt.title(features[j] + ': Non-defaulted')

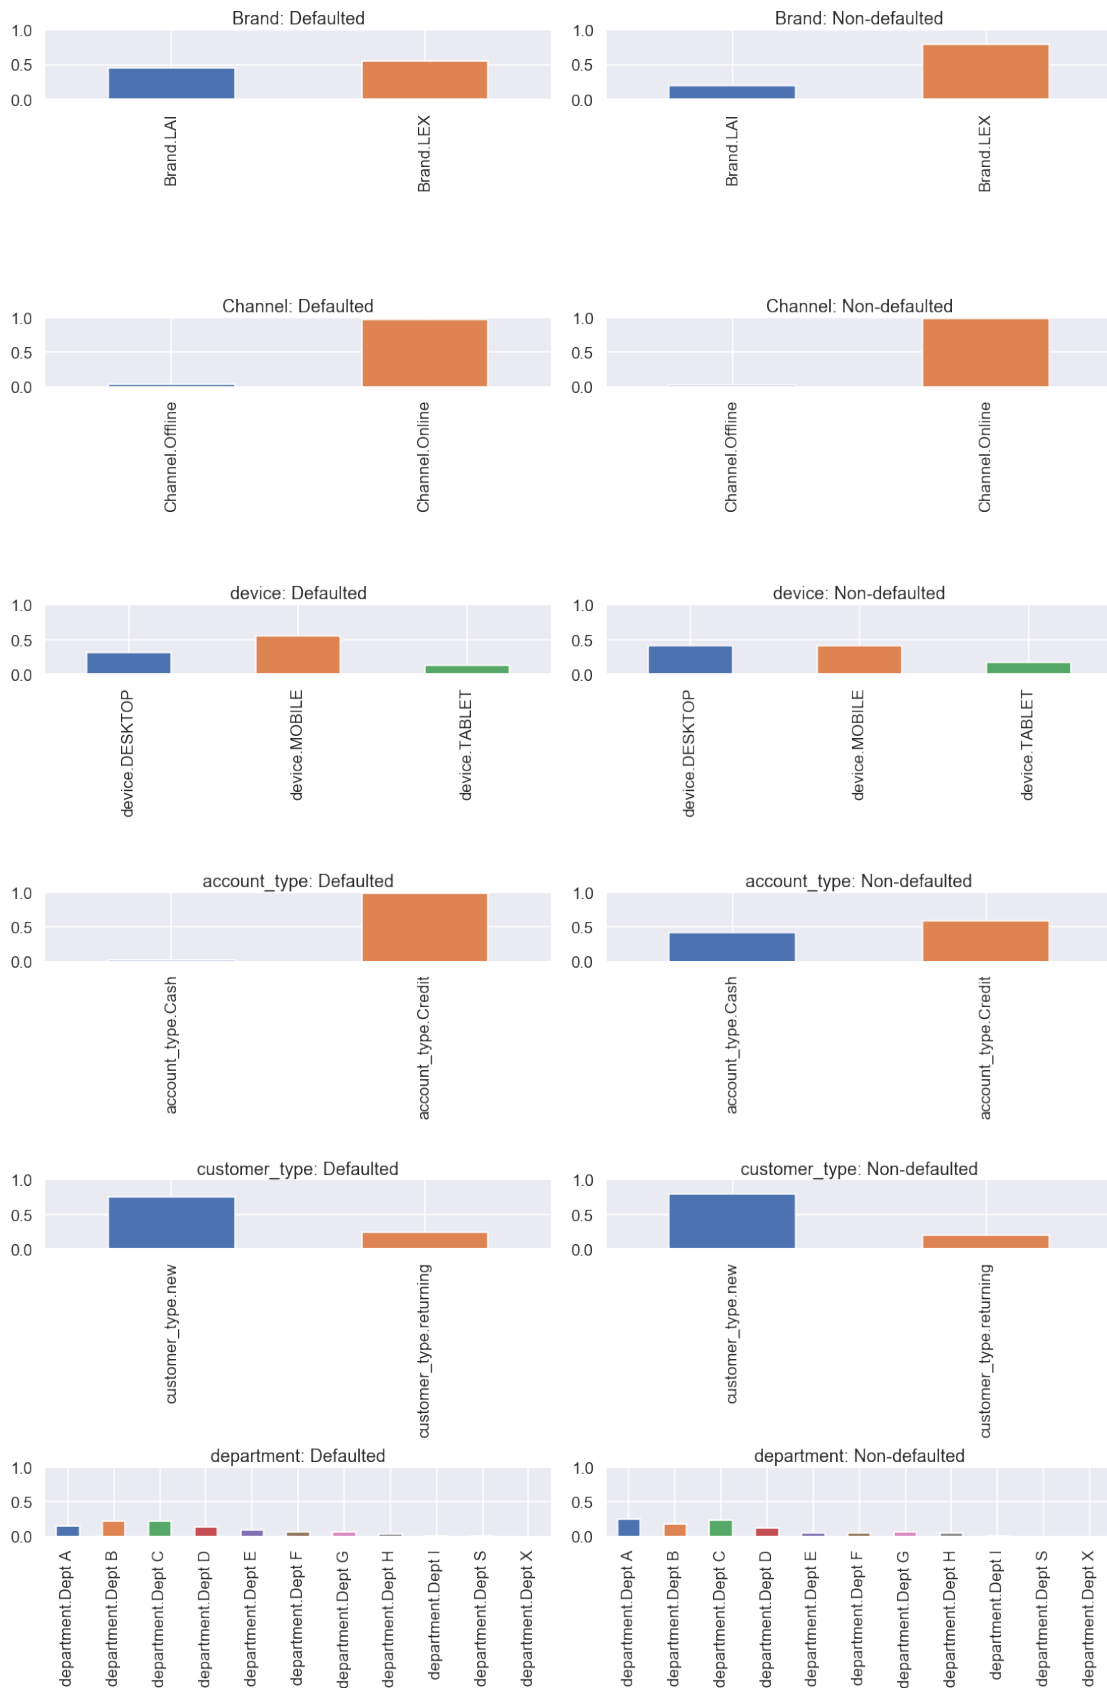
plt.tight_layout()

```

```

In [25]: features = ['Brand', 'Channel', 'device', 'account_type', 'customer_type', 'department']
plot_prob_categorical_unstacked_features_given_defaulted(dfCustomer_2, features)

```



Observations:

- Brand: There is a difference in the distribution of this feature for defaulted and non-defaulted users. There is higher probability of observing Littlewoods for defaulted than for non-defaulted customers.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Channel: There is no visible difference in the distribution of this feature for defaulted and non-defaulted users.
Thus, this feature could be excluded from any model that tries to predict the chance of a user defaulting.
- Device: There is a difference in the distribution of this feature for defaulted and non-defaulted users. There is higher probability of observing Mobile for defaulted than for non-defaulted customers.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Account Type: There is a difference in the distribution of this feature for defaulted and non-defaulted users. There is lower probability of observing Cash for defaulted than for non-defaulted customers.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.
- Customer Type: There is no visible difference in the distribution of this feature for defaulted and non-defaulted users.
Thus, this feature could be excluded from any model that tries to predict the chance of a user defaulting.
- Department: There is a slight difference in the distribution of this feature for defaulted and non-defaulted users. There is higher probability of observing DepartmentB (and lower probability for Department A) for defaulted than for non-defaulted customers.
Thus, this feature should be included in any model that tries to predict the chance of a user defaulting.

Select only relevant features and one-hot-encode the selected categorical features

```
In [26]: with timer('Removing un-informative columns:'):
        selected_columns_orders = []

        selected_numeric_columns_orders = [
            'Gross_Demand_Pre_Credit.min',
            'Gross_Demand_Pre_Credit.max',
            'Gross_Demand_Pre_Credit.mean',
            'days_between_last_and_first_order',
            'number_of_orders',
            'avg_days_between_orders'
        ]
        selected_columns_orders.extend(selected_numeric_columns_orders)
```

```

selected_categorical_columns_orders = [
    'Brand.LAI',
    'Brand.LEX',
    'device.DESKTOP',
    'device.MOBILE',
    'device.TABLET',
    'account_type.Cash',
    'account_type.Credit',
    'department.Dept A',
    'department.Dept B',
    'department.Dept C',
    'department.Dept D',
    'department.Dept E',
    'department.Dept F',
    'department.Dept G',
    'department.Dept H',
    'department.Dept I',
    'department.Dept S',
    'department.Dept X',
]
selected_columns_orders.extend(selected_categorical_columns_orders)

dfCustomer_2 = copy.deepcopy(dfCustomer_2[selected_columns_orders])

dfCustomer_2.head(5)

```

STARTING: Removing un-informative columns:
 FINISHED in 0.00 min.

```

Out[26]:
      Gross_Demand_Pre_Credit.min  Gross_Demand_Pre_Credit.max  \
identifier
1                7.00                110.00
2               79.99                79.99
3               59.99                199.00
4               79.99                79.99
5                9.90                409.00

      Gross_Demand_Pre_Credit.mean  days_between_last_and_first_order  \
identifier
1                38.137097                532
2                79.990000                0
3               129.495000                70
4                79.990000                0
5               210.100000               175

```


	number_of_orders	avg_days_between_orders	Brand.LAI	Brand.LEX	\
identifier					
1	31	17.161290	0	31	
2	1	0.000000	0	1	
3	2	35.000000	0	2	
4	1	0.000000	0	1	
5	9	19.444444	0	9	

	device.DESKTOP	device.MOBILE	...	\
identifier			...	
1	0	27	...	
2	0	1	...	
3	0	2	...	
4	1	0	...	
5	1	3	...	

	department.Dept B	department.Dept C	department.Dept D	\
identifier				
1	1	0	6	
2	0	1	0	
3	0	1	0	
4	0	1	0	
5	1	4	0	

	department.Dept E	department.Dept F	department.Dept G	\
identifier				
1	18	5	1	
2	0	0	0	
3	0	0	0	
4	0	0	0	
5	4	0	0	

	department.Dept H	department.Dept I	department.Dept S	\
identifier				
1	0	0	0	
2	0	0	0	
3	1	0	0	
4	0	0	0	
5	0	0	0	

	department.Dept X
identifier	
1	0
2	0
3	0
4	0
5	0

[5 rows x 24 columns]

2.4.3 Join the two tables into a single data table and log-transform some columns:

This table will be the main table for clustering and classification studies. Each row will be indexed by customer id and each column will represent a feature.

```
In [27]: with timer('Merging the dfCustomer_1 and dfCustomer_2 tables'):
    # merge teables
    dfCustomerFeatures = dfCustomer_1.merge(dfCustomer_2, how='inner', left_index=True)
    # print some stats
    table_stats(dfCustomer_1, 'dfCustomer_1')
    print()
    table_stats(dfCustomer_2, 'dfCustomer_2')
    print()
    table_stats(dfCustomerFeatures, 'dfCustomerFeatures')

with timer('Log-transforming some numeric features'):
    columns_to_log_transform = [
        'LTIME_NET_SALES_AMT',
        'LTIME_NO_ORDERS',
        'days_between_last_payment_and_start',
        'Gross_Demand_Pre_Credit.min',
        'Gross_Demand_Pre_Credit.max',
        'Gross_Demand_Pre_Credit.mean',
        'days_between_last_and_first_order',
        'number_of_orders',
        'avg_days_between_orders'
    ]
    for clmn in columns_to_log_transform:
        # find rows with value equal to zero or less
        idxZero = dfCustomerFeatures[dfCustomerFeatures[clmn] <= 0].index
        # set them to a small value
        dfCustomerFeatures.loc[idxZero, clmn] = 1e-6
        # log transform
        dfCustomerFeatures[clmn] = dfCustomerFeatures[clmn].apply(np.log10)

    #free memory
    del dfCustomer_1, dfCustomer_2
    gc.collect();

    dfCustomerFeatures.head(5)
```

```
STARTING: Merging the dfCustomer_1 and dfCustomer_2 tables
TABLE DFCUSTOMER_1
```

```
-----
Number of rows                               : 509,684
```

Number of columns : 42

Column name Number of unique values (data type)

```
        identifier: 509,684 (int64)
        defaulted: 2 (bool)
    LTIME_NET_SALES_AMT: 101,239 (float64)
    LTIME_NO_ORDERS: 167 (int64)
        birth_year: 96 (float64)
    days_between_last_payment_and_start: 1,173 (float64)
        year_start: 3 (int64)
    credit_band.U1: 2 (int32)
    credit_band.U2: 2 (int32)
    credit_band.U3: 2 (int32)
    credit_band.W1: 2 (int32)
    credit_band.W4: 2 (int32)
    credit_band.W7: 2 (int32)
    credit_band.X1: 2 (int32)
    credit_band.X2: 2 (int32)
    credit_band.X3: 2 (int32)
    credit_band.X4: 2 (int32)
    credit_band.X5: 2 (int32)
    credit_band.Y1: 2 (int32)
    credit_band.Y2: 2 (int32)
    credit_band.Y3: 2 (int32)
    credit_band.Y4: 2 (int32)
    credit_band.Y5: 2 (int32)
    credit_band.Y6: 2 (int32)
    credit_band.Y7: 2 (int32)
    credit_band.Z1: 2 (int32)
    credit_band.Z2: 2 (int32)
    credit_band.Z3: 2 (int32)
    credit_band.Z4: 2 (int32)
    credit_band.Z5: 2 (int32)
    credit_band.Z6: 2 (int32)
    credit_band.Z7: 2 (int32)
    title.2ND-LIEUTENANT: 2 (int32)
        title.DOCTOR: 2 (int32)
            title.FR: 2 (int32)
            title.LORD: 2 (int32)
            title.MISS: 2 (int32)
            title.MR: 2 (int32)
            title.MRS: 2 (int32)
            title.MS: 2 (int32)
        title.PROFESSOR: 2 (int32)
        title.REVEREND: 2 (int32)
```

TABLE DFCUSTOMER_2

```

Number of rows          : 367,669
Number of columns       : 24

```

```

Column name              Number of unique values (data type)

```

```

    Gross_Demand_Pre_Credit.min: 4,158 (float64)
    Gross_Demand_Pre_Credit.max: 5,444 (float64)
    Gross_Demand_Pre_Credit.mean: 127,629 (float64)
    days_between_last_and_first_order: 157 (int64)
        number_of_orders: 330 (int64)
    avg_days_between_orders: 9,320 (float64)
        Brand.LAI: 243 (int32)
        Brand.LEX: 320 (int32)
    device.DESKTOP: 222 (int32)
    device.MOBILE: 256 (int32)
    device.TABLET: 190 (int32)
    account_type.Cash: 102 (int32)
    account_type.Credit: 329 (int32)
    department.Dept A: 184 (int32)
    department.Dept B: 124 (int32)
    department.Dept C: 63 (int32)
    department.Dept D: 85 (int32)
    department.Dept E: 83 (int32)
    department.Dept F: 85 (int32)
    department.Dept G: 67 (int32)
    department.Dept H: 25 (int32)
    department.Dept I: 15 (int32)
    department.Dept S: 10 (int32)
    department.Dept X: 6 (int32)

```

TABLE DFCUSTOMERFEATURES

```

Number of rows          : 367,669
Number of columns       : 66

```

```

Column name              Number of unique values (data type)

```

```

    identifier: 367,669 (int64)
    defaulted: 2 (bool)
    LTIME_NET_SALES_AMT: 99,632 (float64)
    LTIME_NO_ORDERS: 166 (int64)
    birth_year: 95 (float64)
    days_between_last_payment_and_start: 1,173 (float64)
    year_start: 3 (int64)
    credit_band.U1: 2 (int32)

```

```

        credit_band.U2: 2 (int32)
        credit_band.U3: 2 (int32)
        credit_band.W1: 2 (int32)
        credit_band.W4: 1 (int32)
        credit_band.W7: 2 (int32)
        credit_band.X1: 2 (int32)
        credit_band.X2: 2 (int32)
        credit_band.X3: 2 (int32)
        credit_band.X4: 2 (int32)
        credit_band.X5: 2 (int32)
        credit_band.Y1: 2 (int32)
        credit_band.Y2: 2 (int32)
        credit_band.Y3: 2 (int32)
        credit_band.Y4: 2 (int32)
        credit_band.Y5: 2 (int32)
        credit_band.Y6: 2 (int32)
        credit_band.Y7: 2 (int32)
        credit_band.Z1: 2 (int32)
        credit_band.Z2: 2 (int32)
        credit_band.Z3: 2 (int32)
        credit_band.Z4: 2 (int32)
        credit_band.Z5: 2 (int32)
        credit_band.Z6: 2 (int32)
        credit_band.Z7: 2 (int32)
    title.2ND-LIEUTENANT: 1 (int32)
        title.DOCTOR: 2 (int32)
            title.FR: 2 (int32)
            title.LORD: 2 (int32)
            title.MISS: 2 (int32)
            title.MR: 2 (int32)
            title.MRS: 2 (int32)
            title.MS: 2 (int32)
        title.PROFESSOR: 2 (int32)
        title.REVEREND: 1 (int32)
    Gross_Demand_Pre_Credit.min: 4,158 (float64)
    Gross_Demand_Pre_Credit.max: 5,444 (float64)
    Gross_Demand_Pre_Credit.mean: 127,629 (float64)
    days_between_last_and_first_order: 157 (int64)
        number_of_orders: 330 (int64)
    avg_days_between_orders: 9,320 (float64)
        Brand.LAI: 243 (int32)
        Brand.LEX: 320 (int32)
    device.DESKTOP: 222 (int32)
    device.MOBILE: 256 (int32)
    device.TABLET: 190 (int32)
    account_type.Cash: 102 (int32)
    account_type.Credit: 329 (int32)
    department.Dept A: 184 (int32)

```

```

department.Dept B: 124 (int32)
department.Dept C: 63 (int32)
department.Dept D: 85 (int32)
department.Dept E: 83 (int32)
department.Dept F: 85 (int32)
department.Dept G: 67 (int32)
department.Dept H: 25 (int32)
department.Dept I: 15 (int32)
department.Dept S: 10 (int32)
department.Dept X: 6 (int32)

```

FINISHED in 0.01 min.

STARTING: Log-transforming some numeric features
 FINISHED in 0.01 min.

```

Out[27]:
      identifier  defaulted  LTIME_NET_SALES_AMT  LTIME_NO_ORDERS  \
      identifier
37295          37295      False          2.781339          0.698970
315441         315441      False          1.536432          0.000000
489894         489894      False          2.710067          0.698970
419603         419603      False          1.414639          0.000000
399884         399884      False          2.472917          0.778151

      birth_year  days_between_last_payment_and_start  year_start  \
      identifier
37295          1972.0          2.149219          2016
315441         1966.0          1.146128          2016
489894         1996.0          2.296665          2016
419603         1946.0          0.000000          2016
399884         1996.0          NaN          2016

      credit_band.U1  credit_band.U2  credit_band.U3  ...  \
      identifier
37295              0              0              0  ...
315441              0              0              0  ...
489894              0              0              0  ...
419603              0              0              0  ...
399884              0              0              0  ...

      department.Dept B  department.Dept C  department.Dept D  \
      identifier
37295                  1                  1                  0
315441                 0                  0                  0
489894                 2                  2                  0
419603                 0                  1                  0

```

399884	0	2	0
--------	---	---	---

	department.Dept E	department.Dept F	department.Dept G \
identifier			
37295	0	0	1
315441	0	0	0
489894	0	0	0
419603	0	0	0
399884	2	3	0

	department.Dept H	department.Dept I	department.Dept S \
identifier			
37295	0	0	0
315441	0	0	0
489894	0	0	0
419603	0	0	0
399884	0	0	0

	department.Dept X
identifier	
37295	0
315441	0
489894	0
419603	0
399884	0

[5 rows x 66 columns]

2.5 Customer clustering using unsupervised methods

```
In [ ]: import scipy as sp
import sklearn as sk
from sklearn.preprocessing import StandardScaler, MaxAbsScaler, MinMaxScaler, RobustScaler
from sklearn import cluster, mixture, linear_model
from sklearn.cluster import AgglomerativeClustering, DBSCAN
from sklearn.metrics import pairwise_distances
from sklearn.metrics.pairwise import euclidean_distances, cosine_distances
from sklearn import (manifold, datasets, decomposition, ensemble,
                      discriminant_analysis, random_projection, metrics)
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import kneighbors_graph
#from sklearn.feature_selection import f_regression, mutual_info_regression
from sklearn.model_selection import train_test_split, cross_val_score
from scipy.spatial.distance import pdist, squareform
#from scipy.optimize import curve_fit
#from scipy.stats import power_divergence
#from scipy.special import xlogy
import umap
```

2.5.1 Visual inspection of possible clusters using embeddings to lower dimensions

1. I'll embed a sample of the original samples with 64 features into 3D space using three different embedding methods: PCA, Umap, and t-SNE. I'll use different metrics, as well.
2. Then I'll look at the projections of the samples in these lower dimensional spaces to get an idea of how the data (customers) cluster.
3. Then I'll run a few clustering methods in the original data space to cluster the sampled set of data
4. I'll plot the cluster results in various embedded sub-spaces

```
In [28]: class ClusterAnalysis():
    def __init__(self, df_data):
        self.arr_data = df_data.fillna(0).values
        self.dat_column_labels = list(df_data.columns)
        self.embedding_models = {}
        self.embedded_data = {}

    def standardise_data(self, method):
        if method.lower() == 'standard':
            return StandardScaler().fit_transform(self.arr_data)
        elif method.lower() == 'minmax':
            return MinMaxScaler().fit_transform(self.arr_data)
        elif method.lower() == 'maxabs':
            return MaxAbsScaler().fit_transform(self.arr_data)
        elif method.lower() == 'robust':
            return RobustScaler().fit_transform(self.arr_data)
        else:
            return self.arr_data

    def embed_pca(self, n_components=3, algorithm = 'arpack', rescaling_method=None):
        self.embedding_models['pca'] = decomposition.TruncatedSVD(
            n_components = n_components,
            algorithm = algorithm
        )
        self.embedded_data['pca'] = self.embedding_models['pca'].fit_transform(
            self.standardise_data(rescaling_method))

    def embed_umap(self, n_components=3, n_neighbors=5, min_dist=.2, metrics=None, rescaling_method=None):
        if metrics is None:
            self.embedding_models['umap'] = umap.UMAP(
                n_components = n_components,
                n_neighbors = n_neighbors,
                min_dist = min_dist)
            self.embedded_data['umap'] = self.embedding_models['umap'].fit_transform(
                self.standardise_data(rescaling_method))
        else:
            self.embedding_models['umap'] = {}
            self.embedded_data['umap'] = {}
            for metric in metrics:
```



```

        self.embedding_models['umap'][metric] = umap.UMAP(
            n_components = n_components,
            n_neighbors = n_neighbors,
            min_dist = min_dist,
            metric = metric)
        self.embedded_data['umap'][metric] = self.embedding_models['umap'][metric].fit_transform(
            self.standardise_data(rescaling_method))

def embed_tsne(self, n_components=3, perplexity=30, learning_rate=4, early_exaggeration=12,
               metrics=None, rescaling_method=None):
    if metrics is None:
        self.embedding_models['tsne'] = manifold.TSNE(
            n_components = n_components,
            perplexity = perplexity,
            learning_rate = learning_rate,
            early_exaggeration = early_exaggeration,
            n_iter = n_iter,
            init=init)
        self.embedded_data['tsne'] = self.embedding_models['tsne'].fit_transform(
            self.standardise_data(self.arr_data, rescaling_method))
    else:
        self.embedding_models['tsne'] = {}
        self.embedded_data['tsne'] = {}
        for metric in metrics:
            self.embedding_models['tsne'][metric] = manifold.TSNE(
                n_components = n_components,
                perplexity = perplexity,
                learning_rate = learning_rate,
                early_exaggeration = early_exaggeration,
                n_iter = n_iter,
                init=init,
                metric = metric
            )
            self.embedded_data['tsne'][metric] = self.embedding_models['tsne'][metric].fit_transform(
                self.standardise_data(rescaling_method))

def plot_embeddings(self, alpha=0.25, color_labels=None):
    if color_labels is not None:
        n_colors = len(np.unique(color_labels))
    cnt_rows = 0
    for model, item in self.embedded_data.items():
        if model == 'pca':
            cnt_rows += 1
        else:
            #cnt_rows += len(item.keys())
            for metric, item_2 in item.items():
                cnt_rows += 1

```

```

plt.close('all')
plt.figure(figsize=(24, 5 * cnt_rows))
j = 0
for model, item in self.embedded_data.items():
    if model == 'pca':
        for k in range(item.shape[1]):
            for l in range(k+1, item.shape[1]):
                j+=1
                plt.subplot(cnt_rows, 3, j)
                if color_labels is None:
                    plt.plot(item[:, k], item[:, l], '.', alpha = alpha)
                else:
                    for color_ in range(n_colors):
                        idx_color = np.where(color_labels == color_)[0]
                        plt.plot(item[idx_color, k], item[idx_color, l], '.',
                                plt.legend()
                    plt.xticks([], plt.yticks([])
                    plt.title(model + ': P_' + str(k)+str(l))
    else:
        #cnt_rows += len(item.keys())
        for metric, item_2 in item.items():
            for k in range(item_2.shape[1]):
                for l in range(k+1, item_2.shape[1]):
                    j+=1
                    plt.subplot(cnt_rows, 3, j)
                    if color_labels is None:
                        plt.plot(item_2[:, k], item_2[:, l], '.', alpha = alpha)
                    else:
                        for color_ in range(n_colors):
                            idx_color = np.where(color_labels == color_)[0]
                            plt.plot(item_2[idx_color, k], item_2[idx_color, l],
                                    #plt.legend()
                        plt.xticks([], plt.yticks([])
                        plt.title(model + ': ' + metric)

def cluster_data(self, params, rescaling_method=None):

    # =====
    # Create cluster objects
    # =====

    mini_batch_kmeans = cluster.MinibatchKMeans(n_clusters=params['n_clusters'])

    affinity_propagation = cluster.AffinityPropagation(
        damping=params['affinityPropagation_damping'], preference=params['affinity

```

```

# estimate bandwidth for mean shift
bandwidth = cluster.estimate_bandwidth(self.standardise_data(rescaling_method),
mean_shift = cluster.MeanShift(bandwidth=bandwidth, bin_seeding=False)

spectral = cluster.SpectralClustering(
    n_clusters=params['n_clusters'], eigen_solver='arpack',
    affinity="nearest_neighbors")

# connectivity matrix for structured Ward
connectivity = kneighbors_graph(
    self.standardise_data(rescaling_method), n_neighbors=params['ward_n_neigh
# make connectivity symmetric
connectivity = 0.5 * (connectivity + connectivity.T)
ward = cluster.AgglomerativeClustering(
    n_clusters=params['n_clusters'], linkage='ward',
    connectivity=connectivity)

average_linkage = cluster.AgglomerativeClustering(
    linkage="average", affinity="cosine",
    n_clusters=params['n_clusters'], connectivity=connectivity)

dbscan = cluster.DBSCAN(eps=params['dbscan_eps'], metric=params['dbscan_metric'])

birch = cluster.Birch(threshold=params['birch_thrshld'], n_clusters=params['n_clusters'])

gmm = mixture.GaussianMixture(
    n_components=params['n_clusters'], covariance_type='full')

self.clustering_algorithms = [
    ('MiniBatchKMeans', mini_batch_kmeans),
    #('AffinityPropagation', affinity_propagation),
    #('MeanShift', mean_shift),
    ('SpectralClustering', spectral),
    ('Ward', ward),
    ('AgglomerativeClustering', average_linkage),
    #('DBSCAN', dbscan),
    ('Birch', birch),
    ('GaussianMixture', gmm)
]

self.cluster_labels=np.zeros((self.standardise_data(rescaling_method).shape[0],

for j, (name, algorithm) in enumerate(self.clustering_algorithms):
    t0 = tm.time()

    # catch warnings related to kneighbors_graph
    with warnings.catch_warnings():
        warnings.filterwarnings(

```

```

        "ignore",
        message="the number of connected components of the " +
        "connectivity matrix is [0-9]{1,2}" +
        " > 1. Completing it to avoid stopping the tree early.",
        category=UserWarning)
warnings.filterwarnings(
    "ignore",
    message="Graph is not fully connected, spectral embedding" +
    " may not work as expected.",
    category=UserWarning)
algorithm.fit(self.standardise_data(rescaling_method))

if hasattr(algorithm, 'labels_'):
    self.cluster_labels[:,j] = algorithm.labels_.flatten()
    numberClusters = len(np.unique(algorithm.labels_))
else:
    self.cluster_labels[:,j] = algorithm.predict(self.standardise_data(rescaling_method))
    numberClusters = len(algorithm.weights_)

t1 = tm.time()
print('{0:2d}. {1:30s}: time {2:3.3f}; number of clusters = {3:}'.format(
    i, name, t1 - t0, numberClusters))

def plot_clusters(self, data, clustering_algorithms, labels, alpha = 0.01, type_='c'):
    #palette = ["#2ecc71", "#3498db", "#e74c3c", "#9b59b6", "#34495e", ]
    palette = itertools.cycle(sns.color_palette())

    plt.close('all')

    plt.figure(figsize = (24, 32))
    for n, (name, algorithm) in enumerate(clustering_algorithms):
        temp = copy.deepcopy(data)
        #print('Clustering method is', name)
        #temp['labels'] = labels[:,n] #[str(c) for c in labels[:,n]]
        uniqueLabels = np.unique(labels[:,n])
        #print('unique labels', uniqueLabels)

        # Number of clusters in labels, ignoring noise if present.
        n_clusters_ = len(set(labels[:,n])) - (1 if -1 in labels[:,n] else 0)
        #print('Estimated number of clusters: %d' % n_clusters_)
        #print('')

        plt.subplot(5,2,n+1)
        plt.title(name + '; Number clusters = ' + str(n_clusters_))
        for j, label in enumerate(uniqueLabels):

```

```

# print(temp[np.where(labels[:,n] == label)[0],:].shape)
color_ = next(palette)
if type_ == 'features':
    '''
    plt.plot(
        np.transpose(temp[np.where(labels[:,n] == label)[0],:]),
        color=color_,
        alpha=alpha
    )
    '''

    centroid = np.mean(temp[np.where(labels[:,n] == label)[0],:], axis=1)
    plt.plot(
        centroid,
        color=color_,
        linewidth=3
    )

    if xticks_ is None:
        a=1
    else:
        plt.xticks(range(len(centroid)), xticks_, rotation='vertical')

elif type_ == 'scatter':
    plt.plot(
        temp[np.where(labels[:,n] == label)[0],0],
        temp[np.where(labels[:,n] == label)[0],1],
        '.',
        markersize=30,
        color=color_,
        alpha=alpha
    )

    centroid = np.mean(temp[np.where(labels[:,n] == label)[0],:], axis=1)
    plt.plot(
        centroid[0],
        centroid[1],
        color=color_,
        markersize=200,
    )

plt.tight_layout()
plt.show()

```

Note: here, in order to save time, I use a random subset of the whole data table

```

In [29]: with timer('Embedding data'):
    features = [clmn for clmn in dfCustomerFeatures.columns if clmn not in ['identifi
    metrics = [
        'euclidean',
        'cosine',
        'correlation'
    ]
    n_sample = 20000
    selected_rows = np.random.permutation(len(dfCustomerFeatures))[:n_sample]
    objClusterAnalysis = ClusterAnalysis(dfCustomerFeatures.iloc[selected_rows][features])
    with timer('embedding PCA', n_blanks=4):
        objClusterAnalysis.embed_pca(rescaling_method = 'maxabs')
    with timer('embedding UMAP', n_blanks=4):
        objClusterAnalysis.embed_umap(n_neighbors=100, min_dist=.7, rescaling_method = 'maxabs')
    #with timer('embedding TSNE', n_blanks=4):
    #    objClusterAnalysis.embed_tsne(n_components=3, perplexity=30, learning_rate=4e-4)
    with timer('plotting embedding projections', n_blanks=4):
        objClusterAnalysis.plot_embeddings(color_labels = dfCustomerFeatures.iloc[selected_rows]['customer_id'])

```

STARTING: Embedding data

STARTING: embedding PCA

FINISHED in 0.00 min.

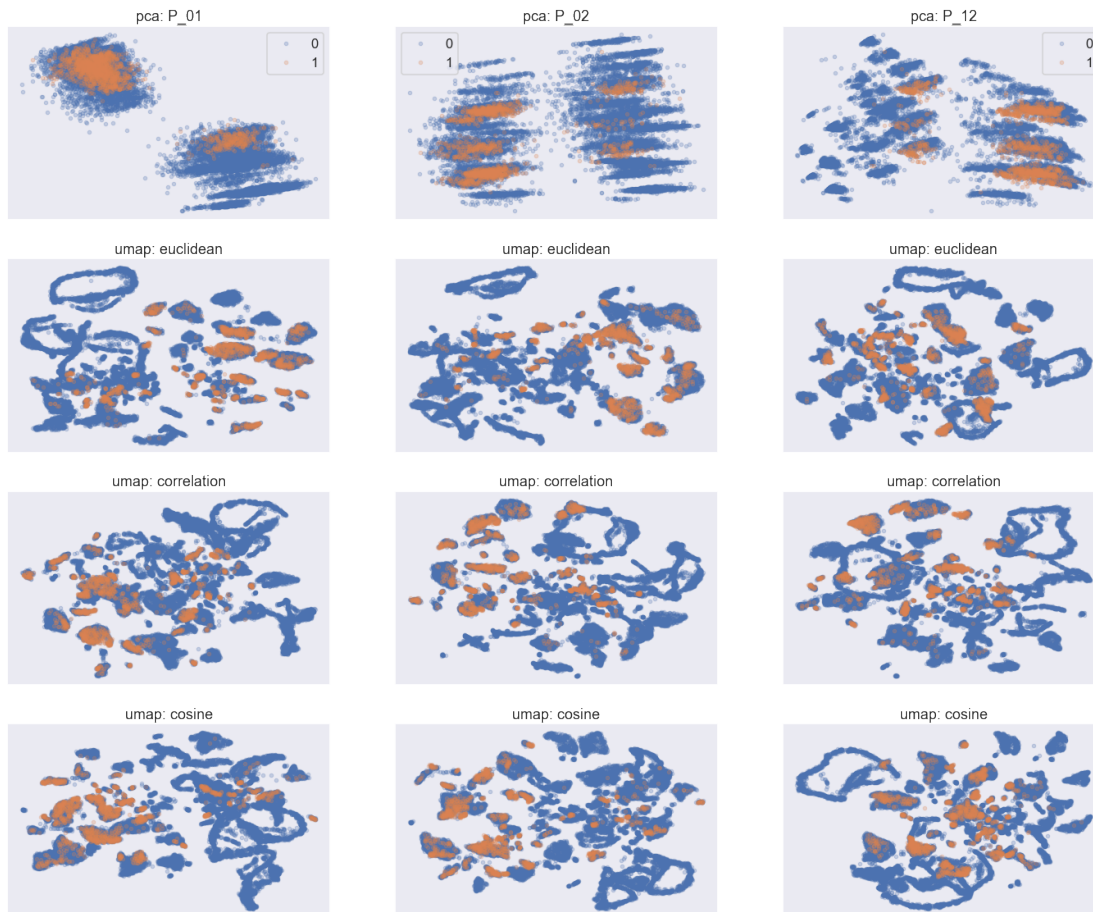
STARTING: embedding UMAP

FINISHED in 2.28 min.

STARTING: plotting embedding projections

FINISHED in 0.00 min.

FINISHED in 2.28 min.



Note: In the above figure, in each panel one sees the projections of the data on two of the directions of the embedding subspace. First column is for directions 1 and 2, second is for 1 and 3, and last is for 2 and 3.

Observations:

- PCA projections: There are two big clusters, which further break down into smaller less distinct ones.
These two big clusters roughly seem to align with whether a customer defaulted or not (orange for defaulted). There are more defaulted customers in one of the clusters.
Thus, we could run a clustering algorithm with two clusters only, then we'd run similar embeddings and clustering for the cluster with more defaulted customers.
- UMAP projections with Euclidean metric: Again, clearly distinguishable two big clusters, which this time break daown into more distinct smaller ones. There seem to be four bigger sub-clusters, two under ech main cluster These two big clusters, too, roughly seem to align with whether a customer defaulted or not (orange for defaulted). There are more defaulted customers in one of the clusters.
Also, the defaulted users are seen only in among distinct sub-clusters.
Furthermore, we can see that some of the smaller clusters for a loop, indicating a continuous

transformations in the original feature space.

Again, we could run a clustering algorithm with two clusters only, then we'd run similar embeddings and clustering for the cluster with more defaulted customers.

- UMAP projections with Cosine metric: Less clearly distinguishable two big clusters, which again break down into distinct smaller ones.
Otherwise, similar to the previous case
- UMAP projections with Correlation metric: Similar to UMAP projections with Correlation metric

2.5.2 Run clustering algorithms

... and colour-code the previous projection by the labels generated by each algorithm.

```
In [30]: with timer('Clustering customers'):
        clusteringParams = {
            'menShift_quantile': .55,
            'affinityPropagation_damping': .999,
            'affinityPropagation_preference': None,
            'affinityPropagation_max_iter': 500,
            'ward_n_neighbors': 10,
            'n_clusters': 2,
            'birch_thrshld': .1,
            'dbscan_eps': 10.01,
            'dbscan_metric': 'euclidean'
        }
        objClusterAnalysis.cluster_data(clusteringParams, rescaling_method = 'maxabs')

        with timer('Plotting cluster results in embeded spaces'):
            with timer('onto PCA projections', n_blanks=4):
                objClusterAnalysis.plot_clusters(objClusterAnalysis.embedded_data['pca'], objC

            with timer('onto UMAP Euclidean projections', n_blanks=4):
                objClusterAnalysis.plot_clusters(objClusterAnalysis.embedded_data['umap']['euc

            with timer('onto UMAP Cosine projections', n_blanks=4):
                objClusterAnalysis.plot_clusters(objClusterAnalysis.embedded_data['umap']['cos
```

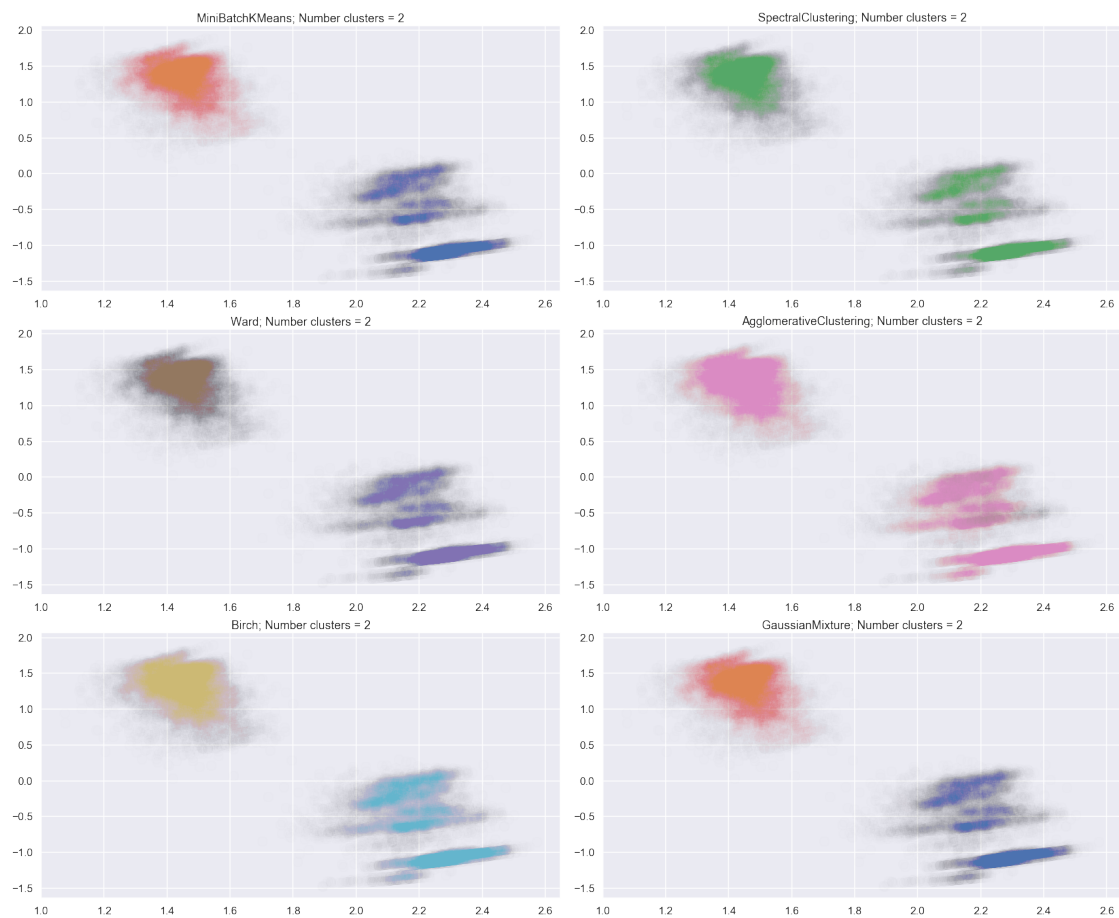
STARTING: Clustering customers

0. MiniBatchKMeans	: time 0.122; number of clusters = 2
1. SpectralClustering	: time 57.233; number of clusters = 2
2. Ward	: time 4.382; number of clusters = 2
3. AgglomerativeClustering	: time 4.543; number of clusters = 2
4. Birch	: time 4.551; number of clusters = 2
5. GaussianMixture	: time 0.432; number of clusters = 2

FINISHED in 1.84 min.

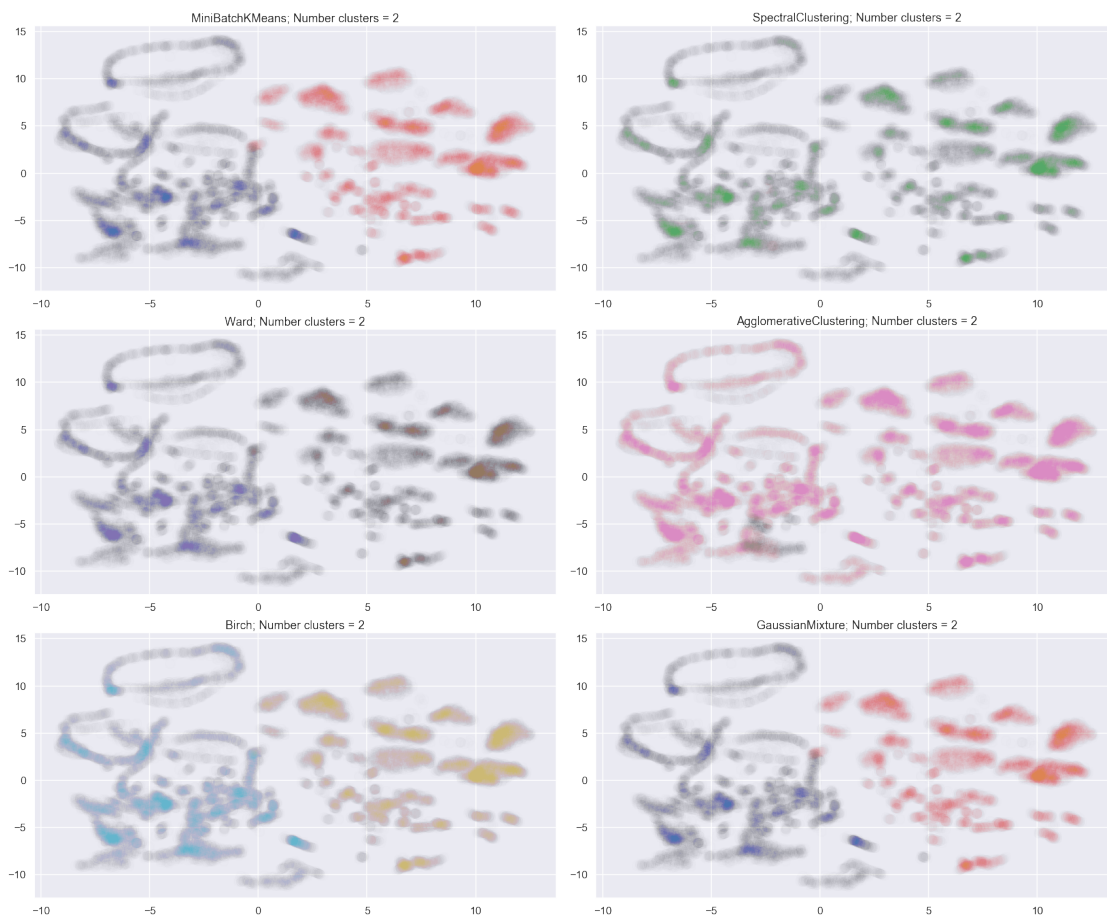
STARTING: Plotting cluster results in embeded spaces

STARTING: onto PCA projections



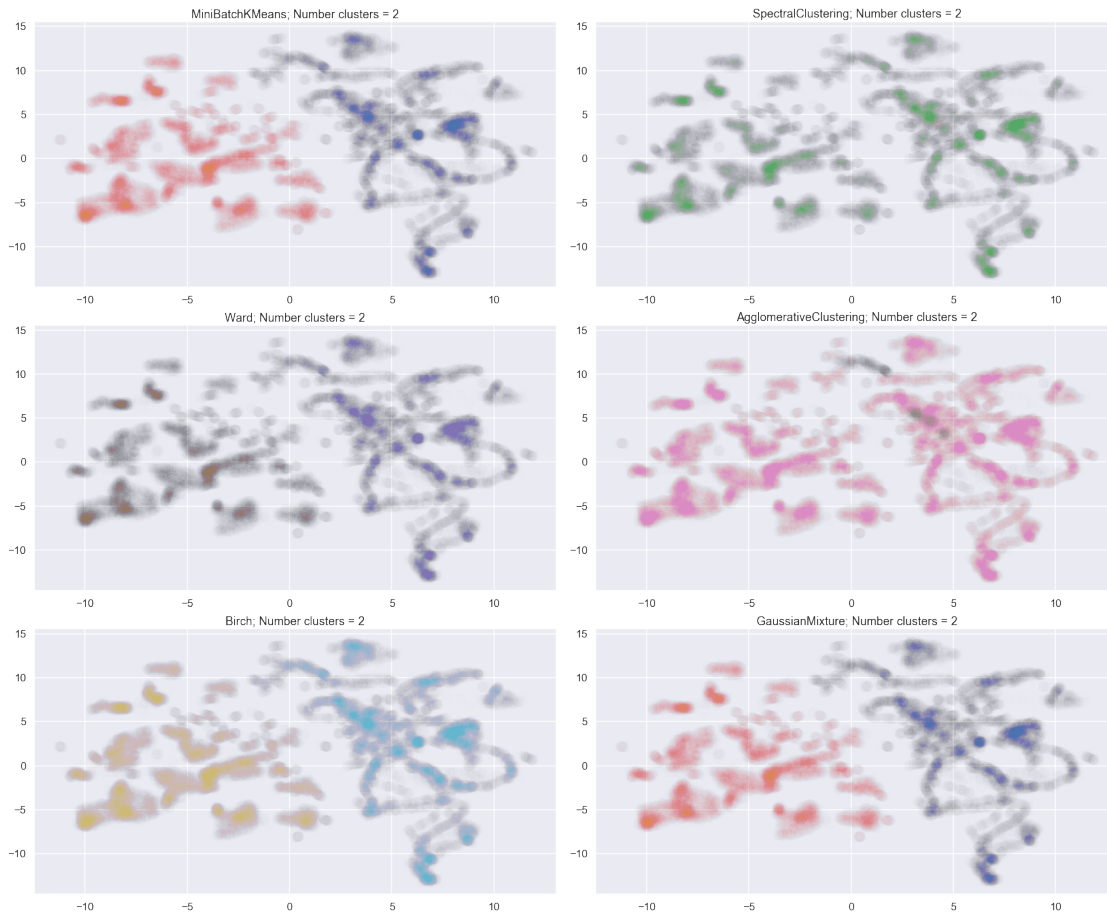
FINISHED in 0.03 min.

STARTING: onto UMAP Euclidean projections



FINISHED in 0.03 min.

STARTING: onto UMAP Cosine projections



FINISHED in 0.03 min.

FINISHED in 0.09 min.

Observations:

- Best algorithms: All algorithms but SpectralClustering seem to produce acceptable results
- Embeddings most consistent with cluster results: All embeddings seem to be doing a reasonable job, with PC clearly standing out and UMP with Euclidean Metric giving a nice balance of separating both bigger clusters and smaller clusters within them

2.5.3 Next steps I'd do, if I had time to do it till the end:

1. Use either MiniBatchKMeans or Ward or GaussianMixture methods results and compare which features of the two clusters differ the most and identify these as important features for identifying customers with high risk of defaults.
2. Select the cluster with more defaulted users.

3. Run sub-space embedding and clustering algorithms on this selected data set.
4. Compare on which features samples in different clusters differ and how.
5. Go to 2.
6. Keep iterating this loop until desired level of granularity has been reached.

2.6 Customer clustering using supervised methods (classification)

```
In [31]: from sklearn import linear_model
        from sklearn import preprocessing
        from sklearn.ensemble import RandomForestRegressor
        from sklearn.model_selection import train_test_split, cross_val_score
        from lightgbm import LGBMRegressor, LGBMClassifier
        from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, accuracy_score
        from sklearn.model_selection import KFold, StratifiedKFold, train_test_split, GridSearchCV
```

2.6.1 Split data into training and validations sets

Note: here, in order to save time, I use a random subset of the whole data table

```
In [32]: with timer('Splitting data to train and validation'):
        n_sample = 50000
        selected_rows = np.random.permutation(len(dfCustomerFeatures))[:n_sample]

        features_classifier = [clmn for clmn in dfCustomerFeatures.columns if clmn not in

        train_x, valid_x, train_y, valid_y = train_test_split(
            dfCustomerFeatures.iloc[selected_rows][features_classifier],
            dfCustomerFeatures.iloc[selected_rows]['defaulted'].astype(int),
            test_size = 0.40,
            random_state= 1984
        )
        print(4*' ' + 'train_x.shape =', train_x.shape)
        print(4*' ' + 'train_y.shape =', train_y.shape)
        print(4*' ' + 'valid_x.shape =', valid_x.shape)
        print(4*' ' + 'valid_y.shape =', valid_y.shape)
```

```
STARTING: Splitting data to train and validation
train_x.shape = (30000, 64)
train_y.shape = (30000,)
valid_x.shape = (20000, 64)
valid_y.shape = (20000,)
FINISHED in 0.00 min.
```

2.6.2 Set initial parameters for LightGBM classifier

```
In [33]: lgbmClassifierParams = {
        'objective': 'xentropy', # 'binary', 'xentlambda
```

```

'metric': 'binary_logloss', #'binary_error', 'xentropy' , 'xentlambda'

'bagging_fraction': 0.75,
'bagging_freq': 10,
'boosting_type': 'gbdt',
#'categorical_feature': [0, 23, 24, 25, 26],
#'early_stopping_rounds': 10,
'feature_fraction': 0.75,
'importance_type': 'gain', # 'split'
'learning_rate': 0.1,
'n_estimators': 512, # 512, 1024,
'num_leaves': 15, #31,
'lambda_l1': 0.01,
'lambda_l2': 5.0,
'verbose': 0,

'min_child_samples': 5,
'min_child_weight': 5.0,
'min_split_gain': 0.5,
'min_data_in_bin': 10,
'min_sum_hessian_in_leaf': 10,
'min_data_in_leaf': 20,
'max_bin': 11,
'xgboost_dart_mode': True,
#'max_depth' : -1,
#'nthread': 4,

#'max_bin': 512,
#'subsample_for_bin': 200,
#'subsample': .25,
#'subsample_freq': 1,
#'colsample_bytree': 0.8,
#'lambda_l1': 0.,
#'lambda_l2': 10.,
#'min_split_gain': 1.0,
#'min_child_weight': 1,
#'min_child_samples': 5,
#'scale_pos_weight': 1,

#'num_boost_round': 20,
'is_unbalance': True,
#'categorical_feature': categorical_features
}

```

2.6.3 Do grid search for some of the parameters for LightGBM classifier

```

In [34]: gridParamsClassifier = {
        'learning_rate': [.05, .1, 1.],

```

```

    #'min_child_samples': [5, 20, 50],
    #'min_child_weight': [1., 5., 10.],
    #'min_split_gain': [.1, .5, 1.],
    'metric': ['binary_logloss', 'binary_error', 'xentropy', 'xentlambda'],
    'importance_type': ['gain', 'split'],
    'n_estimators': [256, 512, 1024],
    #'num_leaves': [7, 9, 15],
    #'min_split_gain': [0.5, 1.0],
    #'max_depth': [1, 5, 10],
    #'feature_fraction': [0.7, 0.8, 0.9],
    #'bagging_fraction': [0.7, 0.8, 0.9],
    #'colsample_bytree': [0.64, 0.65],
    #'subsample': [0.25, 0.50, 0.75],
    #'bagging_freq': [1, 5, 10, 50],
    #'lambda_l1': [0.001, 0.01, .05],
    #'lambda_l2': [0.5, 5.0, 25.],
    #'max_bin': [11, 31, 51],
    #'min_sum_hessian_in_leaf': [0.1, 1.0, 5., 50.],
    #'min_data_in_leaf': [5, 10, 50],
    #'boosting_type': ['gbdt', 'dart']
}

clf = LGBMClassifier(**lgbmClassifierParams)

num_folds= 5
stratified = False

if stratified:
    folds = StratifiedKFold(n_splits= num_folds, shuffle=True, random_state=47)
else:
    folds = KFold(n_splits= num_folds, shuffle=True, random_state=47)

grid = GridSearchCV(clf, gridParamsClassifier, verbose=2, cv=folds, n_jobs=6)
print('grid', grid)

# Run the grid
grid.fit(
    train_x, train_y,
    eval_set=[(train_x, train_y), (valid_x, valid_y)],
    verbose= 4,
    feature_name = features_classifier,
    early_stopping_rounds= 20,
    #categorical_feature= categorical_features
)

grid = GridSearchCV(cv=KFold(n_splits=5, random_state=47, shuffle=True),
    error_score='raise-deprecating',
    estimator=LGBMClassifier(bagging_fraction=0.75, bagging_freq=10, boosting_type='gbdt',

```

```

class_weight=None, colsample_bytree=1.0, feature_fraction=0.75,
importance_type='gain', is_unbalance=True, lambda_l1=0.01,
lambda_l2=5.0, learning_rate=0.1, max_bin=11, max_depth=-1,
...=1.0,
subsample_for_bin=200000, subsample_freq=0, verbose=0,
xgboost_dart_mode=True),
fit_params=None, iid='warn', n_jobs=6,
param_grid={'learning_rate': [0.05, 0.1, 1.0], 'metric': ['binary_logloss', 'binary_err
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=2)
Fitting 5 folds for each of 72 candidates, totalling 360 fits

```

```

[Parallel(n_jobs=6)]: Using backend LokyBackend with 6 concurrent workers.
[Parallel(n_jobs=6)]: Done 29 tasks      | elapsed: 14.5s
[Parallel(n_jobs=6)]: Done 150 tasks    | elapsed: 53.7s
[Parallel(n_jobs=6)]: Done 360 out of 360 | elapsed: 1.9min finished

```

Training until validation scores don't improve for 20 rounds.

[4]	training's binary_logloss: 0.221476	valid_1's binary_logloss: 0.219317
[8]	training's binary_logloss: 0.189428	valid_1's binary_logloss: 0.187547
[12]	training's binary_logloss: 0.165132	valid_1's binary_logloss: 0.163345
[16]	training's binary_logloss: 0.146988	valid_1's binary_logloss: 0.144964
[20]	training's binary_logloss: 0.133848	valid_1's binary_logloss: 0.13195
[24]	training's binary_logloss: 0.125182	valid_1's binary_logloss: 0.123539
[28]	training's binary_logloss: 0.11693	valid_1's binary_logloss: 0.115171
[32]	training's binary_logloss: 0.109081	valid_1's binary_logloss: 0.107495
[36]	training's binary_logloss: 0.103531	valid_1's binary_logloss: 0.102456
[40]	training's binary_logloss: 0.0986961	valid_1's binary_logloss: 0.09783
[44]	training's binary_logloss: 0.0947491	valid_1's binary_logloss: 0.0941831
[48]	training's binary_logloss: 0.0916795	valid_1's binary_logloss: 0.0913966
[52]	training's binary_logloss: 0.0886457	valid_1's binary_logloss: 0.0884472
[56]	training's binary_logloss: 0.0861759	valid_1's binary_logloss: 0.0863032
[60]	training's binary_logloss: 0.0842436	valid_1's binary_logloss: 0.0845307
[64]	training's binary_logloss: 0.0820139	valid_1's binary_logloss: 0.0823364
[68]	training's binary_logloss: 0.0805659	valid_1's binary_logloss: 0.0813648
[72]	training's binary_logloss: 0.078911	valid_1's binary_logloss: 0.0798436
[76]	training's binary_logloss: 0.0777085	valid_1's binary_logloss: 0.0788264
[80]	training's binary_logloss: 0.0765763	valid_1's binary_logloss: 0.077914
[84]	training's binary_logloss: 0.0755249	valid_1's binary_logloss: 0.0771205
[88]	training's binary_logloss: 0.0745966	valid_1's binary_logloss: 0.0765632
[92]	training's binary_logloss: 0.0737024	valid_1's binary_logloss: 0.0758952
[96]	training's binary_logloss: 0.0728745	valid_1's binary_logloss: 0.0752475
[100]	training's binary_logloss: 0.0721177	valid_1's binary_logloss: 0.0747582
[104]	training's binary_logloss: 0.0714687	valid_1's binary_logloss: 0.0743927
[108]	training's binary_logloss: 0.0708465	valid_1's binary_logloss: 0.0739629
[112]	training's binary_logloss: 0.0702505	valid_1's binary_logloss: 0.0735391

[116]	training's binary_logloss: 0.0697876	valid_1's binary_logloss: 0.0733292
[120]	training's binary_logloss: 0.0692635	valid_1's binary_logloss: 0.0729811
[124]	training's binary_logloss: 0.0687147	valid_1's binary_logloss: 0.072688
[128]	training's binary_logloss: 0.0682666	valid_1's binary_logloss: 0.0725064
[132]	training's binary_logloss: 0.0678309	valid_1's binary_logloss: 0.0723409
[136]	training's binary_logloss: 0.0674335	valid_1's binary_logloss: 0.0722154
[140]	training's binary_logloss: 0.0670037	valid_1's binary_logloss: 0.0720368
[144]	training's binary_logloss: 0.0666055	valid_1's binary_logloss: 0.0718031
[148]	training's binary_logloss: 0.0662053	valid_1's binary_logloss: 0.0716191
[152]	training's binary_logloss: 0.0658261	valid_1's binary_logloss: 0.0714225
[156]	training's binary_logloss: 0.0654859	valid_1's binary_logloss: 0.0712934
[160]	training's binary_logloss: 0.0652007	valid_1's binary_logloss: 0.0711845
[164]	training's binary_logloss: 0.0648966	valid_1's binary_logloss: 0.071202
[168]	training's binary_logloss: 0.0645507	valid_1's binary_logloss: 0.0711241
[172]	training's binary_logloss: 0.0642628	valid_1's binary_logloss: 0.0710779
[176]	training's binary_logloss: 0.0639812	valid_1's binary_logloss: 0.0709004
[180]	training's binary_logloss: 0.0636887	valid_1's binary_logloss: 0.0707854
[184]	training's binary_logloss: 0.063409	valid_1's binary_logloss: 0.0707319
[188]	training's binary_logloss: 0.0631403	valid_1's binary_logloss: 0.0707272
[192]	training's binary_logloss: 0.0628833	valid_1's binary_logloss: 0.0707078
[196]	training's binary_logloss: 0.0626216	valid_1's binary_logloss: 0.0706989
[200]	training's binary_logloss: 0.0623916	valid_1's binary_logloss: 0.0707027
[204]	training's binary_logloss: 0.062192	valid_1's binary_logloss: 0.0707155
[208]	training's binary_logloss: 0.0619751	valid_1's binary_logloss: 0.0706382
[212]	training's binary_logloss: 0.0617194	valid_1's binary_logloss: 0.0705728
[216]	training's binary_logloss: 0.0614779	valid_1's binary_logloss: 0.070546
[220]	training's binary_logloss: 0.0612675	valid_1's binary_logloss: 0.0704969
[224]	training's binary_logloss: 0.061036	valid_1's binary_logloss: 0.0704879
[228]	training's binary_logloss: 0.0607817	valid_1's binary_logloss: 0.070482
[232]	training's binary_logloss: 0.0605832	valid_1's binary_logloss: 0.0704828
[236]	training's binary_logloss: 0.0603811	valid_1's binary_logloss: 0.0704251
[240]	training's binary_logloss: 0.0601705	valid_1's binary_logloss: 0.0703963
[244]	training's binary_logloss: 0.0599377	valid_1's binary_logloss: 0.0704188
[248]	training's binary_logloss: 0.0597796	valid_1's binary_logloss: 0.0704207
[252]	training's binary_logloss: 0.0595481	valid_1's binary_logloss: 0.0703991
[256]	training's binary_logloss: 0.0592917	valid_1's binary_logloss: 0.0703368
Did not meet early stopping. Best iteration is:		
[256]	training's binary_logloss: 0.0592917	valid_1's binary_logloss: 0.0703368

```

Out[34]: GridSearchCV(cv=KFold(n_splits=5, random_state=47, shuffle=True),
                      error_score='raise-deprecating',
                      estimator=LGBMClassifier(bagging_fraction=0.75, bagging_freq=10, boosting_type=
                      class_weight=None, colsample_bytree=1.0, feature_fraction=0.75,
                      importance_type='gain', is_unbalance=True, lambda_l1=0.01,
                      lambda_l2=5.0, learning_rate=0.1, max_bin=11, max_depth=-1,
                      ...=1.0,
                      subsample_for_bin=200000, subsample_freq=0, verbose=0,

```



```

        xgboost_dart_mode=True),
    fit_params=None, iid='warn', n_jobs=6,
    param_grid={'learning_rate': [0.05, 0.1, 1.0], 'metric': ['binary_logloss', 'b
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=2)

```

2.6.4 Fit a new classifier with the best parameters from the grid search

```

In [35]: for param,val in grid.best_params_.items():
        lgbmClassifierParams[param] = val
    print('BEST MODEL PARAMETERS:')
    [print(key,':',val) for key,val in lgbmClassifierParams.items()]
    print()
    clf = LGBMClassifier(**lgbmClassifierParams)
    clf.fit(
        train_x, train_y,
        eval_set=[(train_x, train_y), (valid_x, valid_y)],
        verbose= 4,
        feature_name = features_classifier,
        early_stopping_rounds= 20,
        #categorical_feature= categorical_features
    )

    with timer('Saving model', n_blanks=4):
        # save model to file
        clf.booster_.save_model(
            os.path.join(
                PATHS['WORKSPACE'],
                'clf_predict_01.txt'))

        # dump model with pickle
        with open(
            os.path.join(
                PATHS['WORKSPACE'],
                'clf_predict_01.pkl'), 'wb') as fout:
            pickle.dump(clf, fout)

    plt.close('all')
    #plt.plot(clf.evals_result_['training'][lgbmParams['metric']], label='training')
    plt.plot(clf.evals_result_['training'][lgbmClassifierParams['metric']], label='training')
    plt.plot(clf.evals_result_['valid_1'][lgbmClassifierParams['metric']], label='validation')
    plt.legend()

```

```

BEST MODEL PARAMETERS:
lambda_l1 : 0.01
metric : binary_logloss
xgboost_dart_mode : True
min_child_weight : 5.0

```

```

objective : xentropy
importance_type : gain
min_split_gain : 0.5
verbose : 0
min_data_in_bin : 10
bagging_freq : 10
bagging_fraction : 0.75
is_unbalance : True
min_data_in_leaf : 20
min_sum_hessian_in_leaf : 10
learning_rate : 0.05
max_bin : 11
boosting_type : gbdt
n_estimators : 256
feature_fraction : 0.75
min_child_samples : 5
num_leaves : 15
lambda_l2 : 5.0


```

Training until validation scores don't improve for 20 rounds.

[4]	training's binary_logloss: 0.221476	valid_1's binary_logloss: 0.219317
[8]	training's binary_logloss: 0.189428	valid_1's binary_logloss: 0.187547
[12]	training's binary_logloss: 0.165132	valid_1's binary_logloss: 0.163345
[16]	training's binary_logloss: 0.146988	valid_1's binary_logloss: 0.144964
[20]	training's binary_logloss: 0.133848	valid_1's binary_logloss: 0.13195
[24]	training's binary_logloss: 0.125182	valid_1's binary_logloss: 0.123539
[28]	training's binary_logloss: 0.11693	valid_1's binary_logloss: 0.115171
[32]	training's binary_logloss: 0.109081	valid_1's binary_logloss: 0.107495
[36]	training's binary_logloss: 0.103531	valid_1's binary_logloss: 0.102456
[40]	training's binary_logloss: 0.0986961	valid_1's binary_logloss: 0.09783
[44]	training's binary_logloss: 0.0947491	valid_1's binary_logloss: 0.0941831
[48]	training's binary_logloss: 0.0916795	valid_1's binary_logloss: 0.0913966
[52]	training's binary_logloss: 0.0886457	valid_1's binary_logloss: 0.0884472
[56]	training's binary_logloss: 0.0861759	valid_1's binary_logloss: 0.0863032
[60]	training's binary_logloss: 0.0842436	valid_1's binary_logloss: 0.0845307
[64]	training's binary_logloss: 0.0820139	valid_1's binary_logloss: 0.0823364
[68]	training's binary_logloss: 0.0805659	valid_1's binary_logloss: 0.0813648
[72]	training's binary_logloss: 0.078911	valid_1's binary_logloss: 0.0798436
[76]	training's binary_logloss: 0.0777085	valid_1's binary_logloss: 0.0788264
[80]	training's binary_logloss: 0.0765763	valid_1's binary_logloss: 0.077914
[84]	training's binary_logloss: 0.0755249	valid_1's binary_logloss: 0.0771205
[88]	training's binary_logloss: 0.0745966	valid_1's binary_logloss: 0.0765632
[92]	training's binary_logloss: 0.0737024	valid_1's binary_logloss: 0.0758952
[96]	training's binary_logloss: 0.0728745	valid_1's binary_logloss: 0.0752475
[100]	training's binary_logloss: 0.0721177	valid_1's binary_logloss: 0.0747582
[104]	training's binary_logloss: 0.0714687	valid_1's binary_logloss: 0.0743927
[108]	training's binary_logloss: 0.0708465	valid_1's binary_logloss: 0.0739629
[112]	training's binary_logloss: 0.0702505	valid_1's binary_logloss: 0.0735391

[116]	training's binary_logloss: 0.0697876	valid_1's binary_logloss: 0.0733292
[120]	training's binary_logloss: 0.0692635	valid_1's binary_logloss: 0.0729811
[124]	training's binary_logloss: 0.0687147	valid_1's binary_logloss: 0.072688
[128]	training's binary_logloss: 0.0682666	valid_1's binary_logloss: 0.0725064
[132]	training's binary_logloss: 0.0678309	valid_1's binary_logloss: 0.0723409
[136]	training's binary_logloss: 0.0674335	valid_1's binary_logloss: 0.0722154
[140]	training's binary_logloss: 0.0670037	valid_1's binary_logloss: 0.0720368
[144]	training's binary_logloss: 0.0666055	valid_1's binary_logloss: 0.0718031
[148]	training's binary_logloss: 0.0662053	valid_1's binary_logloss: 0.0716191
[152]	training's binary_logloss: 0.0658261	valid_1's binary_logloss: 0.0714225
[156]	training's binary_logloss: 0.0654859	valid_1's binary_logloss: 0.0712934
[160]	training's binary_logloss: 0.0652007	valid_1's binary_logloss: 0.0711845
[164]	training's binary_logloss: 0.0648966	valid_1's binary_logloss: 0.071202
[168]	training's binary_logloss: 0.0645507	valid_1's binary_logloss: 0.0711241
[172]	training's binary_logloss: 0.0642628	valid_1's binary_logloss: 0.0710779
[176]	training's binary_logloss: 0.0639812	valid_1's binary_logloss: 0.0709004
[180]	training's binary_logloss: 0.0636887	valid_1's binary_logloss: 0.0707854
[184]	training's binary_logloss: 0.063409	valid_1's binary_logloss: 0.0707319
[188]	training's binary_logloss: 0.0631403	valid_1's binary_logloss: 0.0707272
[192]	training's binary_logloss: 0.0628833	valid_1's binary_logloss: 0.0707078
[196]	training's binary_logloss: 0.0626216	valid_1's binary_logloss: 0.0706989
[200]	training's binary_logloss: 0.0623916	valid_1's binary_logloss: 0.0707027
[204]	training's binary_logloss: 0.062192	valid_1's binary_logloss: 0.0707155
[208]	training's binary_logloss: 0.0619751	valid_1's binary_logloss: 0.0706382
[212]	training's binary_logloss: 0.0617194	valid_1's binary_logloss: 0.0705728
[216]	training's binary_logloss: 0.0614779	valid_1's binary_logloss: 0.070546
[220]	training's binary_logloss: 0.0612675	valid_1's binary_logloss: 0.0704969
[224]	training's binary_logloss: 0.061036	valid_1's binary_logloss: 0.0704879
[228]	training's binary_logloss: 0.0607817	valid_1's binary_logloss: 0.070482
[232]	training's binary_logloss: 0.0605832	valid_1's binary_logloss: 0.0704828
[236]	training's binary_logloss: 0.0603811	valid_1's binary_logloss: 0.0704251
[240]	training's binary_logloss: 0.0601705	valid_1's binary_logloss: 0.0703963
[244]	training's binary_logloss: 0.0599377	valid_1's binary_logloss: 0.0704188
[248]	training's binary_logloss: 0.0597796	valid_1's binary_logloss: 0.0704207
[252]	training's binary_logloss: 0.0595481	valid_1's binary_logloss: 0.0703991
[256]	training's binary_logloss: 0.0592917	valid_1's binary_logloss: 0.0703368
Did not meet early stopping. Best iteration is:		
[256]	training's binary_logloss: 0.0592917	valid_1's binary_logloss: 0.0703368
STARTING: Saving model		
FINISHED in 0.00 min.		

Out[35]: <matplotlib.legend.Legend at 0x29e822d8eb8>



2.6.5 Display feature importances, based on the trained model


```
In [36]: def display_importances(feature_importance_df_):
    cols = feature_importance_df_[["feature", "importance"]].groupby("feature").mean()
    best_features = feature_importance_df_.loc[feature_importance_df_.feature.isin(cols.index)]
    plt.close('all')
    plt.figure(figsize=(24, 24))
    sns.barplot(x="importance", y="feature", data=best_features.sort_values(by="importance", ascending=False))
    plt.title('LightGBM Features (avg over folds)')
    plt.tight_layout()
    plt.show()

In [37]: if False:
    feature_imp = pd.DataFrame(sorted(zip(clf.feature_importances_, train_x.columns)))

    plt.figure(figsize=(20, 10))
    sns.barplot(x="Value", y="Feature", data=feature_imp.sort_values(by="Value", ascending=False))
    plt.title('LightGBM Features (avg over folds)')
    plt.tight_layout()
    plt.show()
else:
    feature_importance_df = pd.DataFrame()
    feature_importance_df["feature"] = features_classifier
    feature_importance_df["importance"] = clf.feature_importances_
```

```
display_importances(feature_importance_df)
```

```
sumFI = feature_importance_df['importance'].sum()
feature_importance_df['importance'] /= sumFI
feature_importance_df.sort_values(by='importance', ascending=False)[:10]
```



```
Out[37]:
```

	feature	importance
3	days_between_last_payment_and_start	0.227785
7	credit_band.U3	0.137253
52	account_type.Credit	0.134956
4	year_start	0.121871
0	LTIME_NET_SALES_AMT	0.089386
5	credit_band.U1	0.057595

18	credit_band.Y3	0.040513
43	days_between_last_and_first_order	0.029602
16	credit_band.Y1	0.028241
46	Brand.LAI	0.027193

2.6.6 Check model performance

Since this is highly imbalanced data, aka. there are many more customers who do not default than defaulted (roughly 9 to 1), I'll check how well the trained model performs against the null model of predicting only non-default

```
In [38]: y_pred = clf.predict(valid_x)
          print('The accuracy of prediction is:', accuracy_score(valid_y, y_pred))
          print('The roc_auc_score of prediction is:', roc_auc_score(valid_y, y_pred))
          print('The accuracy of always predicting one class is:', max(valid_y.mean(), 1 - val.
```

```
The accuracy of prediction is: 0.9737
The roc_auc_score of prediction is: 0.8836433340559903
The accuracy of always predicting one class is: 0.92035
```

Observation: So, the model performs better than the null model, but it still might need class-balancing the data.

2.6.7 Balance the data set:

I'll check the probabilities of defaulting that the model produces for each sample, then play with the threshold in such a way that when the model predicts non-defaulted it's accurate for almost all of the cases (high precision for non-defaulted) and is only about 5% of the times wrong when it says defaulted (high recall for defaulted).

Thus in the set of predicted defaults, there will be almost all actual defaulted customers, plus only a fraction of all non-defaulted ones.

Then I'll train a second model on this second set alone.

The final system will be composed of two stages:

1. First, filter out the customers for which we are fairly confident they will not default.
2. Next, focus on the remaining ones.

Run the classifier for various thresholds and monitor recall, precision and f1 score

```
In [39]: #=====
          # Confusion Matrix
          # rows: actual
          # cols: predicted
          #=====

          tblPredictions_Valid = pd.DataFrame({'actual': valid_y})
          tblPredictions_Valid['probability'] = clf.predict_proba(valid_x, num_iteration=clf.be
```

```

grid_thrslds = np.linspace(0.001,0.1, 10)
grid_thrslds
for thrshld in grid_thrslds:
    tblPredictions_Valid['predicted'] = (tblPredictions_Valid['probability'] >= thrshld)
    cm = confusion_matrix(tblPredictions_Valid['actual'], tblPredictions_Valid['predicted'])
    recall = np.array([cm[0,0]/np.sum(cm[0,:]), cm[1,1]/np.sum(cm[1,:])])
    precission = np.array([cm[0,0]/np.sum(cm[:,0]), cm[1,1]/np.sum(cm[:,1])])
    f1 = 2 * (precission * recall) / (precission + recall)
    print('\nThreshold =',thrshld)
    print('recall:',recall, np.sqrt(recall[0]*recall[1]))
    print('precission:', precission, np.sqrt(precission[0]*precission[1]))
    print('F1:', f1, np.sqrt(f1[0]*f1[1]))
    print('cm[:,0]/np.sum(cm,axis=1)', cm[:,0]/np.sum(cm,axis=1))
    print('cm[:,0]',cm[:,0])

```

Threshold = 0.001

```

recall: [0.41842777 1.          ] 0.6468599323217078
precission: [1.          0.12953326] 0.35990729006264127
F1: [0.58998813 0.22935714] 0.3678559337916082
cm[:,0]/np.sum(cm,axis=1) [0.41842777 0.          ]
cm[:,0] [7702    0]

```

Threshold = 0.012

```

recall: [0.78774379 0.98618958] 0.8813992965857973
precission: [0.99848506 0.2867835 ] 0.5351159098914741
F1: [0.88068268 0.44435016] 0.6255649386733703
cm[:,0]/np.sum(cm,axis=1) [0.78774379 0.01381042]
cm[:,0] [14500    22]

```

Threshold = 0.023000000000000003

```

recall: [0.85880372 0.9742624 ] 0.9147131614512369
precission: [0.99741309 0.37388581] 0.6106706151516869
F1: [0.92293321 0.54038997] 0.7062179911222087
cm[:,0]/np.sum(cm,axis=1) [0.85880372 0.0257376 ]
cm[:,0] [15808    41]

```

Threshold = 0.034

```

recall: [0.89335579 0.96421846] 0.9281110595580282
precission: [0.99654566 0.43898257] 0.6614122565545911
F1: [0.94213361 0.60329929] 0.7539154726044324
cm[:,0]/np.sum(cm,axis=1) [0.89335579 0.03578154]
cm[:,0] [16444    57]

```

Threshold = 0.045000000000000005

```

recall: [0.91302222 0.95480226] 0.9336785735998363
precission: [0.99573409 0.4871877 ] 0.6964979555939649

```

```
F1: [0.95258608 0.64517497] 0.7839545279970394
cm[:,0]/np.sum(cm,axis=1) [0.91302222 0.04519774]
cm[:,0] [16806    72]
```

```
Threshold = 0.056000000000000001
recall: [0.92660401 0.94978029] 0.9381205805035248
precision: [0.99533147 0.52828212] 0.7251315879274832
F1: [0.95973891 0.67893202] 0.8072158772701432
cm[:,0]/np.sum(cm,axis=1) [0.92660401 0.05021971]
cm[:,0] [17056    80]
```

```
Threshold = 0.067
recall: [0.93681752 0.94224733] 0.939528501017373
precision: [0.99469312 0.56343844] 0.7486309775770266
F1: [0.96488823 0.70519145] 0.824882373011152
cm[:,0]/np.sum(cm,axis=1) [0.93681752 0.05775267]
cm[:,0] [17244    92]
```

```
Threshold = 0.078000000000000001
recall: [0.94474928 0.94161959] 0.9431831347980971
precision: [0.99468055 0.59594756] 0.7699204124375624
F1: [0.96907216 0.72992701] 0.8410421780254836
cm[:,0]/np.sum(cm,axis=1) [0.94474928 0.05838041]
cm[:,0] [17390    93]
```

```
Threshold = 0.089000000000000001
recall: [0.95121421 0.93659761] 0.9438776201838109
precision: [0.99426462 0.62426778] 0.7878371474832838
F1: [0.9722631 0.74918403] 0.8534658681185733
cm[:,0]/np.sum(cm,axis=1) [0.95121421 0.06340239]
cm[:,0] [17509   101]
```

```
Threshold = 0.1
recall: [0.95528875 0.93157564] 0.94335768978518
precision: [0.99383937 0.64325964] 0.7995603547342092
F1: [0.97418283 0.76102564] 0.8610331638448737
cm[:,0]/np.sum(cm,axis=1) [0.95528875 0.06842436]
cm[:,0] [17584   109]
```

Select athreshold for the clasifier that will wilter out the customers for which we are confident they will not default

```
In [43]: THRSILD = 0.090
```

```
#tblVisitorFeatures = copy.deepcopy(tblVisitorFeatures_orgnl)
tblPredictions_Train = pd.DataFrame({'actual': train_y})
tblPredictions_Train['probability'] = clf.predict_proba(train_x, num_iteration=clf.be
```



```

tblPredictions_Train['predicted'] = (tblPredictions_Train['probability'] >= THRESHLD).

tblPredictions_Valid = pd.DataFrame({'actual': valid_y})
tblPredictions_Valid['probability'] = clf.predict_proba(valid_x, num_iteration=clf.best_iteration_)
tblPredictions_Valid['predicted'] = (tblPredictions_Valid['probability'] >= THRESHLD).

cm = confusion_matrix(tblPredictions_Valid['actual'], tblPredictions_Valid['predicted'])
recall = np.array([cm[0,0]/np.sum(cm[0,:]), cm[1,1]/np.sum(cm[1,:])])
precision = np.array([cm[0,0]/np.sum(cm[:,0]), cm[1,1]/np.sum(cm[:,1])])
f1 = 2 * (precision * recall) / (precision + recall)
print('\n\nThreshold =', THRESHLD)
print('recall:', recall, np.sqrt(recall[0]*recall[1]))
print('precision:', precision, np.sqrt(precision[0]*precision[1]))
print('F1:', f1, np.sqrt(f1[0]*f1[1]))
print('cm[:,0]/np.sum(cm,axis=1)', cm[:,0]/np.sum(cm,axis=1))
print('cm[:,0]', cm[:,0])

#Print Confusion Matrix
plt.figure()
labels = ['Non-default', 'Defaulted']
plt.figure(figsize=(8,6))
sns.heatmap(cm, xticklabels = labels, yticklabels = labels, annot = True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.ylabel('True Class')
plt.xlabel('Predicted Class')
plt.show()

plt.figure()
false_positive_rate, recall_, thresholds = roc_curve(tblPredictions_Valid['actual'], tblPredictions_Valid['probability'])
roc_auc = auc(false_positive_rate, recall_)
plt.title('Receiver Operating Characteristic (ROC)')
plt.plot(false_positive_rate, recall_, 'b', label = 'AUC = %0.3f' %roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1], [0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.0])
plt.ylabel('Recall')
plt.xlabel('Fall-out (1-Specificity)')

plt.tight_layout()
plt.show()

```


```


Threshold = 0.09
recall: [0.9515945  0.93659761] 0.9440662798188143
precision: [0.9942669  0.62610155] 0.7889943286614143
F1: [0.9724628  0.75050302] 0.8543045524470365

```

```
cm[:,0]/np.sum(cm,axis=1) [0.9515945  0.06340239]
cm[:,0] [17516   101]
```

<Figure size 432x288 with 0 Axes>





Observation: The set of predicted defaults seems much more balanced now.

2.6.8 Select a new set of training and testing

```
In [44]: train_x_2 = train_x.loc[tblPredictions_Train[tblPredictions_Train['predicted'] == 1].index]
train_y_2 = train_y[tblPredictions_Train[tblPredictions_Train['predicted'] == 1].index]
tblPredictions_Train_2 = tblPredictions_Train[tblPredictions_Train['predicted'] == 1].index

valid_x_2 = valid_x.loc[tblPredictions_Valid[tblPredictions_Valid['predicted'] == 1].index]
valid_y_2 = valid_y[tblPredictions_Valid[tblPredictions_Valid['predicted'] == 1].index]
tblPredictions_Valid_2 = tblPredictions_Valid[tblPredictions_Valid['predicted'] == 1].index

# sanity check
print(len(train_y_2)/len(train_y),
      len(valid_y_2)/len(valid_y))
print(len(train_y_2[tblPredictions_Train_2['actual'] == 1])/len(train_y_2),
      len(valid_y_2[tblPredictions_Valid_2['actual'] == 1])/len(valid_y_2))
```

0.1197 0.11915

0.6427179058758006 0.6261015526647083

2.6.9 New grid search and a second model:

Starting with the parameters for the latest model, do grid search for a new model on the new subset of data

```
In [45]: gridParamsClassifier = {
    'learning_rate': [.05, .1, 1.],
    #'min_child_samples': [5, 20, 50],
    #'min_child_weight': [1., 5., 10.],
    #'min_split_gain': [.1, .5, 1.],
    'metric': ['binary_logloss', 'binary_error', 'xentropy', 'xentlambd'],
    'importance_type': ['gain', 'split'],
    'n_estimators': [256, 512, 1024],
    #'num_leaves': [7, 9, 15],
    #'min_split_gain': [0.5, 1.0],
    #'max_depth': [1, 5, 10],
    #'feature_fraction': [0.7, 0.8, 0.9],
    #'bagging_fraction': [0.7, 0.8, 0.9],
    #'colsample_bytree': [0.64, 0.65],
    #'subsample': [0.25, 0.50, 0.75],
    #'bagging_freq': [1, 5, 10, 50],
    #'lambda_l1': [0.001, 0.01, .05],
    #'lambda_l2': [0.5, 5.0, 25.],
    #'max_bin': [11, 31, 51],
    #'min_sum_hessian_in_leaf': [0.1, 1.0, 5., 50.],
    #'min_data_in_leaf': [5, 10, 50],
    #'boosting_type': ['gbdt', 'dart']
}

clf_2 = LGBMClassifier(**lgbmClassifierParams)

num_folds= 5
stratified = False

if stratified:
    folds = StratifiedKFold(n_splits= num_folds, shuffle=True, random_state=47)
else:
    folds = KFold(n_splits= num_folds, shuffle=True, random_state=47)

grid = GridSearchCV(clf_2, gridParamsClassifier, verbose=2, cv=folds, n_jobs=6)
print('grid', grid)

# Run the grid
grid.fit(
    train_x_2, train_y_2,
    eval_set=[(train_x_2, train_y_2), (valid_x_2, valid_y_2)],
    verbose= 4,
    feature_name = features_classifier,
    early_stopping_rounds= 20,
```

```

        #categorical_feature= categorical_features
    )

    # get the best params
    for param,val in grid.best_params_.items():
        lgbmClassifierParams[param] = val
    print('BEST MODEL PARAMETERS:')
    [print(key,':',val) for key,val in lgbmClassifierParams.items()]
    print()

    #fit a new model with best params
    clf_2 = LGBMClassifier(**lgbmClassifierParams)
    clf_2.fit(
        train_x_2, train_y_2,
        eval_set=[(train_x_2, train_y_2), (valid_x_2, valid_y_2)],
        verbose= 4,
        feature_name = features_classifier,
        early_stopping_rounds= 20,
        #categorical_feature= categorical_features
    )

    with timer('Saving model', n_blanks=4):
        # save model to file
        clf_2.booster_.save_model(
            os.path.join(
                PATHS['WORKSPACE'],
                'clf_predict_02.txt'))

        # dump model with pickle
        with open(
            os.path.join(
                PATHS['WORKSPACE'],
                'clf_predict_02.pkl'), 'wb') as fout:
            pickle.dump(clf_2, fout)

    plt.close('all')
    #plt.plot(clf.evals_result_['training'][lgbmParams['metric']], label='training')
    plt.plot(clf.evals_result_['training'][lgbmClassifierParams['metric']], label='training')
    plt.plot(clf.evals_result_['valid_1'][lgbmClassifierParams['metric']], label='validation')
    plt.legend()

    grid = GridSearchCV(cv=KFold(n_splits=5, random_state=47, shuffle=True),
        error_score='raise-deprecating',
        estimator=LGBMClassifier(bagging_fraction=0.75, bagging_freq=10, boosting_type='gbdt',
            class_weight=None, colsample_bytree=1.0, feature_fraction=0.75,
            importance_type='gain', is_unbalance=True, lambda_l1=0.01,
            lambda_l2=5.0, learning_rate=0.05, max_bin=11, max_depth=-1,
            ...=1.0,

```

```

        subsample_for_bin=200000, subsample_freq=0, verbose=0,
        xgboost_dart_mode=True),
    fit_params=None, iid='warn', n_jobs=6,
    param_grid={'learning_rate': [0.05, 0.1, 1.0], 'metric': ['binary_logloss', 'binary_err
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=2)

```

Fitting 5 folds for each of 72 candidates, totalling 360 fits

```

[Parallel(n_jobs=6)]: Using backend LokyBackend with 6 concurrent workers.
[Parallel(n_jobs=6)]: Done 29 tasks      | elapsed:    2.1s
[Parallel(n_jobs=6)]: Done 150 tasks    | elapsed:    9.1s
[Parallel(n_jobs=6)]: Done 360 out of 360 | elapsed:   19.7s finished

```

Training until validation scores don't improve for 20 rounds.

[4]	training's binary_logloss: 0.609124	valid_1's binary_logloss: 0.622025
[8]	training's binary_logloss: 0.571005	valid_1's binary_logloss: 0.586194
[12]	training's binary_logloss: 0.548262	valid_1's binary_logloss: 0.565728
[16]	training's binary_logloss: 0.522544	valid_1's binary_logloss: 0.540234
[20]	training's binary_logloss: 0.502779	valid_1's binary_logloss: 0.520384
[24]	training's binary_logloss: 0.486471	valid_1's binary_logloss: 0.505532
[28]	training's binary_logloss: 0.474739	valid_1's binary_logloss: 0.495437
[32]	training's binary_logloss: 0.465547	valid_1's binary_logloss: 0.48806
[36]	training's binary_logloss: 0.457019	valid_1's binary_logloss: 0.480825
[40]	training's binary_logloss: 0.447358	valid_1's binary_logloss: 0.472064
[44]	training's binary_logloss: 0.440314	valid_1's binary_logloss: 0.466459
[48]	training's binary_logloss: 0.433494	valid_1's binary_logloss: 0.461038
[52]	training's binary_logloss: 0.426912	valid_1's binary_logloss: 0.455134
[56]	training's binary_logloss: 0.421354	valid_1's binary_logloss: 0.450184
[60]	training's binary_logloss: 0.416655	valid_1's binary_logloss: 0.447033
[64]	training's binary_logloss: 0.411562	valid_1's binary_logloss: 0.441849
[68]	training's binary_logloss: 0.407763	valid_1's binary_logloss: 0.439055
[72]	training's binary_logloss: 0.403477	valid_1's binary_logloss: 0.435599
[76]	training's binary_logloss: 0.399585	valid_1's binary_logloss: 0.433311
[80]	training's binary_logloss: 0.395983	valid_1's binary_logloss: 0.430908
[84]	training's binary_logloss: 0.392614	valid_1's binary_logloss: 0.428403
[88]	training's binary_logloss: 0.389415	valid_1's binary_logloss: 0.426887
[92]	training's binary_logloss: 0.385736	valid_1's binary_logloss: 0.424258
[96]	training's binary_logloss: 0.383164	valid_1's binary_logloss: 0.422798
[100]	training's binary_logloss: 0.380511	valid_1's binary_logloss: 0.420911
[104]	training's binary_logloss: 0.377965	valid_1's binary_logloss: 0.419965
[108]	training's binary_logloss: 0.375412	valid_1's binary_logloss: 0.419165
[112]	training's binary_logloss: 0.372933	valid_1's binary_logloss: 0.418482
[116]	training's binary_logloss: 0.370387	valid_1's binary_logloss: 0.418143
[120]	training's binary_logloss: 0.367902	valid_1's binary_logloss: 0.417114
[124]	training's binary_logloss: 0.36565	valid_1's binary_logloss: 0.416619
[128]	training's binary_logloss: 0.363587	valid_1's binary_logloss: 0.416487

[132]	training's binary_logloss: 0.361586	valid_1's binary_logloss: 0.415255
[136]	training's binary_logloss: 0.359654	valid_1's binary_logloss: 0.414588
[140]	training's binary_logloss: 0.35782	valid_1's binary_logloss: 0.414268
[144]	training's binary_logloss: 0.356173	valid_1's binary_logloss: 0.414285
[148]	training's binary_logloss: 0.354341	valid_1's binary_logloss: 0.414508
[152]	training's binary_logloss: 0.352514	valid_1's binary_logloss: 0.414358
[156]	training's binary_logloss: 0.350901	valid_1's binary_logloss: 0.41492
[160]	training's binary_logloss: 0.349565	valid_1's binary_logloss: 0.415574
[164]	training's binary_logloss: 0.347905	valid_1's binary_logloss: 0.415367

Early stopping, best iteration is:

[146]	training's binary_logloss: 0.355175	valid_1's binary_logloss: 0.414061
-------	-------------------------------------	------------------------------------

BEST MODEL PARAMETERS:

```

lambda_l1 : 0.01
metric : binary_logloss
xgboost_dart_mode : True
min_child_weight : 5.0
objective : xentropy
importance_type : gain
min_split_gain : 0.5
verbose : 0
min_data_in_bin : 10
bagging_freq : 10
bagging_fraction : 0.75
is_unbalance : True
min_data_in_leaf : 20
min_sum_hessian_in_leaf : 10
learning_rate : 0.05
max_bin : 11
boosting_type : gbdt
n_estimators : 256
feature_fraction : 0.75
min_child_samples : 5
num_leaves : 15
lambda_l2 : 5.0

```

Training until validation scores don't improve for 20 rounds.

[4]	training's binary_logloss: 0.609124	valid_1's binary_logloss: 0.622025
[8]	training's binary_logloss: 0.571005	valid_1's binary_logloss: 0.586194
[12]	training's binary_logloss: 0.548262	valid_1's binary_logloss: 0.565728
[16]	training's binary_logloss: 0.522544	valid_1's binary_logloss: 0.540234
[20]	training's binary_logloss: 0.502779	valid_1's binary_logloss: 0.520384
[24]	training's binary_logloss: 0.486471	valid_1's binary_logloss: 0.505532
[28]	training's binary_logloss: 0.474739	valid_1's binary_logloss: 0.495437
[32]	training's binary_logloss: 0.465547	valid_1's binary_logloss: 0.48806
[36]	training's binary_logloss: 0.457019	valid_1's binary_logloss: 0.480825
[40]	training's binary_logloss: 0.447358	valid_1's binary_logloss: 0.472064
[44]	training's binary_logloss: 0.440314	valid_1's binary_logloss: 0.466459
[48]	training's binary_logloss: 0.433494	valid_1's binary_logloss: 0.461038


[52]	training's binary_logloss: 0.426912	valid_1's binary_logloss: 0.455134
[56]	training's binary_logloss: 0.421354	valid_1's binary_logloss: 0.450184
[60]	training's binary_logloss: 0.416655	valid_1's binary_logloss: 0.447033
[64]	training's binary_logloss: 0.411562	valid_1's binary_logloss: 0.441849
[68]	training's binary_logloss: 0.407763	valid_1's binary_logloss: 0.439055
[72]	training's binary_logloss: 0.403477	valid_1's binary_logloss: 0.435599
[76]	training's binary_logloss: 0.399585	valid_1's binary_logloss: 0.433311
[80]	training's binary_logloss: 0.395983	valid_1's binary_logloss: 0.430908
[84]	training's binary_logloss: 0.392614	valid_1's binary_logloss: 0.428403
[88]	training's binary_logloss: 0.389415	valid_1's binary_logloss: 0.426887
[92]	training's binary_logloss: 0.385736	valid_1's binary_logloss: 0.424258
[96]	training's binary_logloss: 0.383164	valid_1's binary_logloss: 0.422798
[100]	training's binary_logloss: 0.380511	valid_1's binary_logloss: 0.420911
[104]	training's binary_logloss: 0.377965	valid_1's binary_logloss: 0.419965
[108]	training's binary_logloss: 0.375412	valid_1's binary_logloss: 0.419165
[112]	training's binary_logloss: 0.372933	valid_1's binary_logloss: 0.418482
[116]	training's binary_logloss: 0.370387	valid_1's binary_logloss: 0.418143
[120]	training's binary_logloss: 0.367902	valid_1's binary_logloss: 0.417114
[124]	training's binary_logloss: 0.36565	valid_1's binary_logloss: 0.416619
[128]	training's binary_logloss: 0.363587	valid_1's binary_logloss: 0.416487
[132]	training's binary_logloss: 0.361586	valid_1's binary_logloss: 0.415255
[136]	training's binary_logloss: 0.359654	valid_1's binary_logloss: 0.414588
[140]	training's binary_logloss: 0.35782	valid_1's binary_logloss: 0.414268
[144]	training's binary_logloss: 0.356173	valid_1's binary_logloss: 0.414285
[148]	training's binary_logloss: 0.354341	valid_1's binary_logloss: 0.414508
[152]	training's binary_logloss: 0.352514	valid_1's binary_logloss: 0.414358
[156]	training's binary_logloss: 0.350901	valid_1's binary_logloss: 0.41492
[160]	training's binary_logloss: 0.349565	valid_1's binary_logloss: 0.415574
[164]	training's binary_logloss: 0.347905	valid_1's binary_logloss: 0.415367

Early stopping, best iteration is:

[146]	training's binary_logloss: 0.355175	valid_1's binary_logloss: 0.414061
-------	-------------------------------------	------------------------------------

STARTING: Saving model
FINISHED in 0.00 min.

Out[45]: <matplotlib.legend.Legend at 0x29e9d79c438>



2.6.10 Check model performance

```
In [46]: y_pred_2 = clf_2.predict(valid_x_2)
         print('The accuracy of prediction is:', accuracy_score(valid_y_2, y_pred_2))
         print('The roc_auc_score of prediction is:', roc_auc_score(valid_y_2, y_pred_2))
         print('The accuracy of always predicting one class is:',
               max(valid_y_2.mean(), 1 - valid_y_2.mean()))
```

The accuracy of prediction is: 0.8107427612253462

The roc_auc_score of prediction is: 0.793479176633779

The accuracy of always predicting one class is: 0.6261015526647083

Observation: The fitted model vastly outperforms random guess or guessing only one class.

2.6.11 Select best threshold:

Use f1 score to select best threshold for the new classifier

```
In [47]: tblPredictions_Valid = pd.DataFrame({'actual': valid_y_2})
         tblPredictions_Valid['probability'] = clf_2.predict_proba(valid_x_2, num_iteration=cl

         grid_thrlds = np.linspace(0.3, 0.6, 10)
         grid_thrlds
```

```

for thrshld in grid_thrshlds:
    tblPredictions_Valid['predicted'] = (tblPredictions_Valid['probability'] >= thrshld)
    cm = confusion_matrix(tblPredictions_Valid['actual'], tblPredictions_Valid['predicted'])
    recall = np.array([cm[0,0]/np.sum(cm[0,:]), cm[1,1]/np.sum(cm[1,:])])
    precission = np.array([cm[0,0]/np.sum(cm[:,0]), cm[1,1]/np.sum(cm[:,1])])
    f1 = 2 * (precission * recall) / (precission + recall)
    print('\nThreshold =',thrshld)
    print('recall:',recall, np.sqrt(recall[0]*recall[1]))
    print('precission:', precission, np.sqrt(precission[0]*precission[1]))
    print('F1:', f1, np.sqrt(f1[0]*f1[1]))
    print('cm[:,0]/np.sum(cm,axis=1)', cm[:,0]/np.sum(cm,axis=1))
    print('cm[:,0]',cm[:,0])

```

Threshold = 0.3

```

recall: [0.41301908 0.95174263] 0.6269671953492263
precission: [0.83636364 0.73082862] 0.7818174201526366
F1: [0.55296769 0.82678311] 0.6761540890153248
cm[:,0]/np.sum(cm,axis=1) [0.41301908 0.04825737]
cm[:,0] [368 72]

```

Threshold = 0.3333333333333333

```

recall: [0.47923681 0.93632708] 0.6698674527569357
precission: [0.81800766 0.75067168] 0.7836167354393573
F1: [0.60438783 0.83328363] 0.7096664573037669
cm[:,0]/np.sum(cm,axis=1) [0.47923681 0.06367292]
cm[:,0] [427 95]

```

Threshold = 0.36666666666666664

```

recall: [0.5375982 0.92359249] 0.7046429350169472
precission: [0.80775717 0.7698324 ] 0.7885668267519315
F1: [0.64555256 0.83973187] 0.7362683338691935
cm[:,0]/np.sum(cm,axis=1) [0.5375982 0.07640751]
cm[:,0] [479 114]

```

Threshold = 0.4

```

recall: [0.60381594 0.91353887] 0.7427040671245592
precission: [0.8065967 0.79428904] 0.8004192172103435
F1: [0.69062901 0.84975062] 0.7660694701749375
cm[:,0]/np.sum(cm,axis=1) [0.60381594 0.08646113]
cm[:,0] [538 129]

```

Threshold = 0.43333333333333335

```

recall: [0.65544332 0.89477212] 0.7658148663539445
precission: [0.78812416 0.81303289] 0.8004816413213243
F1: [0.71568627 0.85194639] 0.7808497557653901
cm[:,0]/np.sum(cm,axis=1) [0.65544332 0.10522788]
cm[:,0] [584 157]

```

```

Threshold = 0.4666666666666667
recall: [0.69584736 0.87734584] 0.7813442211412118
precision: [0.77210461 0.82848101] 0.7997962286000203
F1: [0.73199528 0.85221354] 0.78982041495443
cm[:,0]/np.sum(cm,axis=1) [0.69584736 0.12265416]
cm[:,0] [620 183]

```

```

Threshold = 0.5
recall: [0.72502806 0.86193029] 0.7905211244227947
precision: [0.75821596 0.83997387] 0.7980486192950116
F1: [0.74125072 0.85081045] 0.7941434748149198
cm[:,0]/np.sum(cm,axis=1) [0.72502806 0.13806971]
cm[:,0] [646 206]

```

```

Threshold = 0.5333333333333333
recall: [0.7687991 0.84316354] 0.8051232027690756
precision: [0.74537541 0.85928962] 0.8003082838937287
F1: [0.75690608 0.8511502 ] 0.8026460997035546
cm[:,0]/np.sum(cm,axis=1) [0.7687991 0.15683646]
cm[:,0] [685 234]

```

```

Threshold = 0.5666666666666667
recall: [0.79349046 0.82372654] 0.8084671622905186
precision: [0.72886598 0.86978061] 0.7962119662245034
F1: [0.75980656 0.84612737] 0.8018061611357316
cm[:,0]/np.sum(cm,axis=1) [0.79349046 0.17627346]
cm[:,0] [707 263]

```

```

Threshold = 0.6
recall: [0.81144781 0.80764075] 0.8095420431129163
precision: [0.71584158 0.8776402 ] 0.7926230843878502
F1: [0.76065229 0.84118674] 0.7999066295270267
cm[:,0]/np.sum(cm,axis=1) [0.81144781 0.19235925]
cm[:,0] [723 287]

```

2.6.12 Observation:

Best threshold seems to be about 0.45

In [48]: THRESHLD = 0.45

```

tblPredictions_Valid = pd.DataFrame({'actual': valid_y_2})
tblPredictions_Valid['probability'] = clf_2.predict_proba(valid_x_2, num_iteration=cl
tblPredictions_Valid['predicted'] = (tblPredictions_Valid['probability'] >= THRESHLD).

```

```

cm = confusion_matrix(tblPredictions_Valid['actual'], tblPredictions_Valid['predicted'])
recall = np.array([cm[0,0]/np.sum(cm[0,:]), cm[1,1]/np.sum(cm[1,:])])
precision = np.array([cm[0,0]/np.sum(cm[:,0]), cm[1,1]/np.sum(cm[:,1])])
f1 = 2 * (precision * recall) / (precision + recall)
print('\n\nThreshold =', THRESHLD)
print('recall:', recall, np.sqrt(recall[0]*recall[1]))
print('precision:', precision, np.sqrt(precision[0]*precision[1]))
print('F1:', f1, np.sqrt(f1[0]*f1[1]))
print('cm[:,0]/np.sum(cm,axis=1)', cm[:,0]/np.sum(cm,axis=1))
print('cm[:,0]', cm[:,0])

#Print Confusion Matrix
plt.figure()
labels = ['Non-default', 'Defaulted']
plt.figure(figsize=(8,6))
sns.heatmap(cm, xticklabels = labels, yticklabels = labels, annot = True, fmt='d', cm
plt.title('Confusion Matrix')
plt.ylabel('True Class')
plt.xlabel('Predicted Class')
plt.show()

plt.figure()
false_positive_rate, recall_, thresholds = roc_curve(tblPredictions_Valid['actual'],
roc_auc = auc(false_positive_rate, recall_)
plt.title('Receiver Operating Characteristic (ROC)')
plt.plot(false_positive_rate, recall_, 'b', label = 'AUC = %0.3f' %roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1], [0,1], 'r--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.0])
plt.ylabel('Recall')
plt.xlabel('Fall-out (1-Specificity)')

plt.tight_layout()
plt.show()


```


```

Threshold = 0.45
recall: [0.67003367 0.88806971] 0.7713861573490546
precision: [0.78141361 0.81840642] 0.7996961423594078
F1: [0.72145015 0.85181614] 0.7839278539508419
cm[:,0]/np.sum(cm,axis=1) [0.67003367 0.11193029]
cm[:,0] [597 167]

```

<Figure size 432x288 with 0 Axes>






```
In [49]: if False:
    feature_imp = pd.DataFrame(sorted(zip(clf_2.feature_importances_, train_x_2.columns)),
                                columns=['importance', 'feature'])

    plt.figure(figsize=(20, 10))
    sns.barplot(x="Value", y="Feature", data=feature_imp.sort_values(by="Value", ascending=False))
    plt.title('LightGBM Features (avg over folds)')
    plt.tight_layout()
    plt.show()

else:
    feature_importance_df = pd.DataFrame()
    feature_importance_df["feature"] = features_classifier
    feature_importance_df["importance"] = clf_2.feature_importances_
    display_importances(feature_importance_df)

    sumFI = feature_importance_df['importance'].sum()
    feature_importance_df['importance'] /= sumFI
    feature_importance_df.sort_values(by='importance', ascending=False)[:10]
```



```

Out [49]:
          feature  importance
3  days_between_last_payment_and_start  0.310259
4                year_start            0.269377
43  days_between_last_and_first_order    0.047922
0                LTIME_NET_SALES_AMT    0.043315
45      avg_days_between_orders        0.034964
44      number_of_orders            0.031725
16      credit_band.Y1             0.025332
18      credit_band.Y3             0.020360
52      account_type.Credit        0.019094
19      credit_band.Y4             0.017047

```

2.6.13 Observation:

While before balancing the dataset `days_between_last_payment_and_start` and `days_between_last_payment_and_start` were among the most important features, now they have become even more prominent in detecting customers that might default.

2.7 Best and worst customers:

Best customers are those that are least likely to default and generate the most revenue, which I'll measure by the total sales a customer generates (`LTIME_NET_SALES_AMT`).

Customer value will be the product of the two measures: $(1 - \text{Prob}(\text{default}|\text{features})) \times \text{TotalRevenue}$.

I'll define worst customers as those that are most likely to default and generate the least revenue

2.7.1 First, get predictions for all customers with the first model (catcher of non-defaulting customers):

```
In [50]: tblPredictions = pd.DataFrame({'actual': dfCustomerFeatures['defaulted'].astype(int),
tblPredictions['probability'] = clf.predict_proba(dfCustomerFeatures[features], num_i
tblPredictions['total_revenue'] = dfCustomerFeatures['LTIME_NET_SALES_AMT']
tblPredictions['customer_value'] = (1 - tblPredictions['probability']) * tblPrediction
```

2.7.2 Top 20 customers

```
In [51]: tblPredictions.sort_values(by='customer_value', ascending=False).head(20)
```

```
Out [51]:
```

	actual	probability	total_revenue	customer_value
identifier				
328199	0	0.012648	4.346431	4.291459
465063	0	0.038217	4.452349	4.282196
283271	0	0.011285	4.183307	4.136099
104343	0	0.008523	4.135119	4.099875
88199	0	0.025435	4.203445	4.096532
485939	0	0.016994	4.140219	4.069859
83159	0	0.012525	4.116779	4.065217
300620	0	0.007279	4.094956	4.065148
372080	0	0.020329	4.145769	4.061491
236829	0	0.007644	4.091630	4.060355
170510	0	0.012361	4.108292	4.057512
455546	0	0.009072	4.091845	4.054725
308941	0	0.026766	4.162677	4.051257
121609	0	0.007867	4.072497	4.040457
430161	0	0.009522	4.075609	4.036801
153083	0	0.028240	4.143778	4.026758
153568	0	0.014764	4.076085	4.015907
36042	0	0.007238	4.043469	4.014203
455060	0	0.009251	4.050987	4.013512
327394	0	0.006695	4.037221	4.010190

2.7.3 Worst 20 customers

```
In [52]: # get predictions with the second model (catcher of defaulting customers)
tblPredictions_2 = pd.DataFrame({'actual': dfCustomerFeatures['defaulted'].astype(int),
tblPredictions_2['probability'] = clf_2.predict_proba(dfCustomerFeatures[features], n
tblPredictions_2['total_revenue'] = dfCustomerFeatures['LTIME_NET_SALES_AMT']
tblPredictions_2['customer_value'] = (1 - tblPredictions_2['probability']) * tblPredi

tblPredictions[tblPredictions['total_revenue'] > 0].sort_values(by='customer_value', a
```

```
Out [52]:
```

	actual	probability	total_revenue	customer_value
identifier				
485117	1	0.999124	2.355068	0.002064
493358	1	0.999050	2.173186	0.002064
275914	1	0.998958	2.250420	0.002346
253145	1	0.998695	2.146128	0.002800
357838	1	0.998620	2.159868	0.002980
329070	1	0.998947	2.836324	0.002986
460496	1	0.998622	2.191311	0.003020
180837	1	0.998594	2.158362	0.003036
471896	1	0.998537	2.079181	0.003041
503873	1	0.998531	2.110590	0.003101
236365	1	0.998705	2.395588	0.003103
432087	1	0.998705	2.403721	0.003112
298055	1	0.998559	2.170262	0.003127
168883	1	0.998703	2.420038	0.003140
364392	1	0.998699	2.438701	0.003174
487431	1	0.998771	2.597146	0.003191
378973	1	0.998686	2.443106	0.003210
375378	1	0.998484	2.154424	0.003267
272550	1	0.998679	2.492411	0.003292
380258	1	0.998733	2.618100	0.003316